

# Concurrency Control

## Concurrent transactions

When two or more transactions appear to be running against a database at the same time.

However: The CPU can execute only one instruction at a time.

Transactions are **interleaved**, meaning that the operating system quickly switches CPU services among tasks so that some portion of each task is carried out in a given interval, but the transactions appear to be concurrent.

# Concurrency Control

## Concurrency Defined:

Multiple transactions can execute at the same time without compromising data integrity

# Concurrency Control

## Concurrency challenges

**Lost Update:** One user's changes overlay or interfere with another user's changes.

**Inconsistent Reads:** Two users can read the same thing at the same time and get different results.

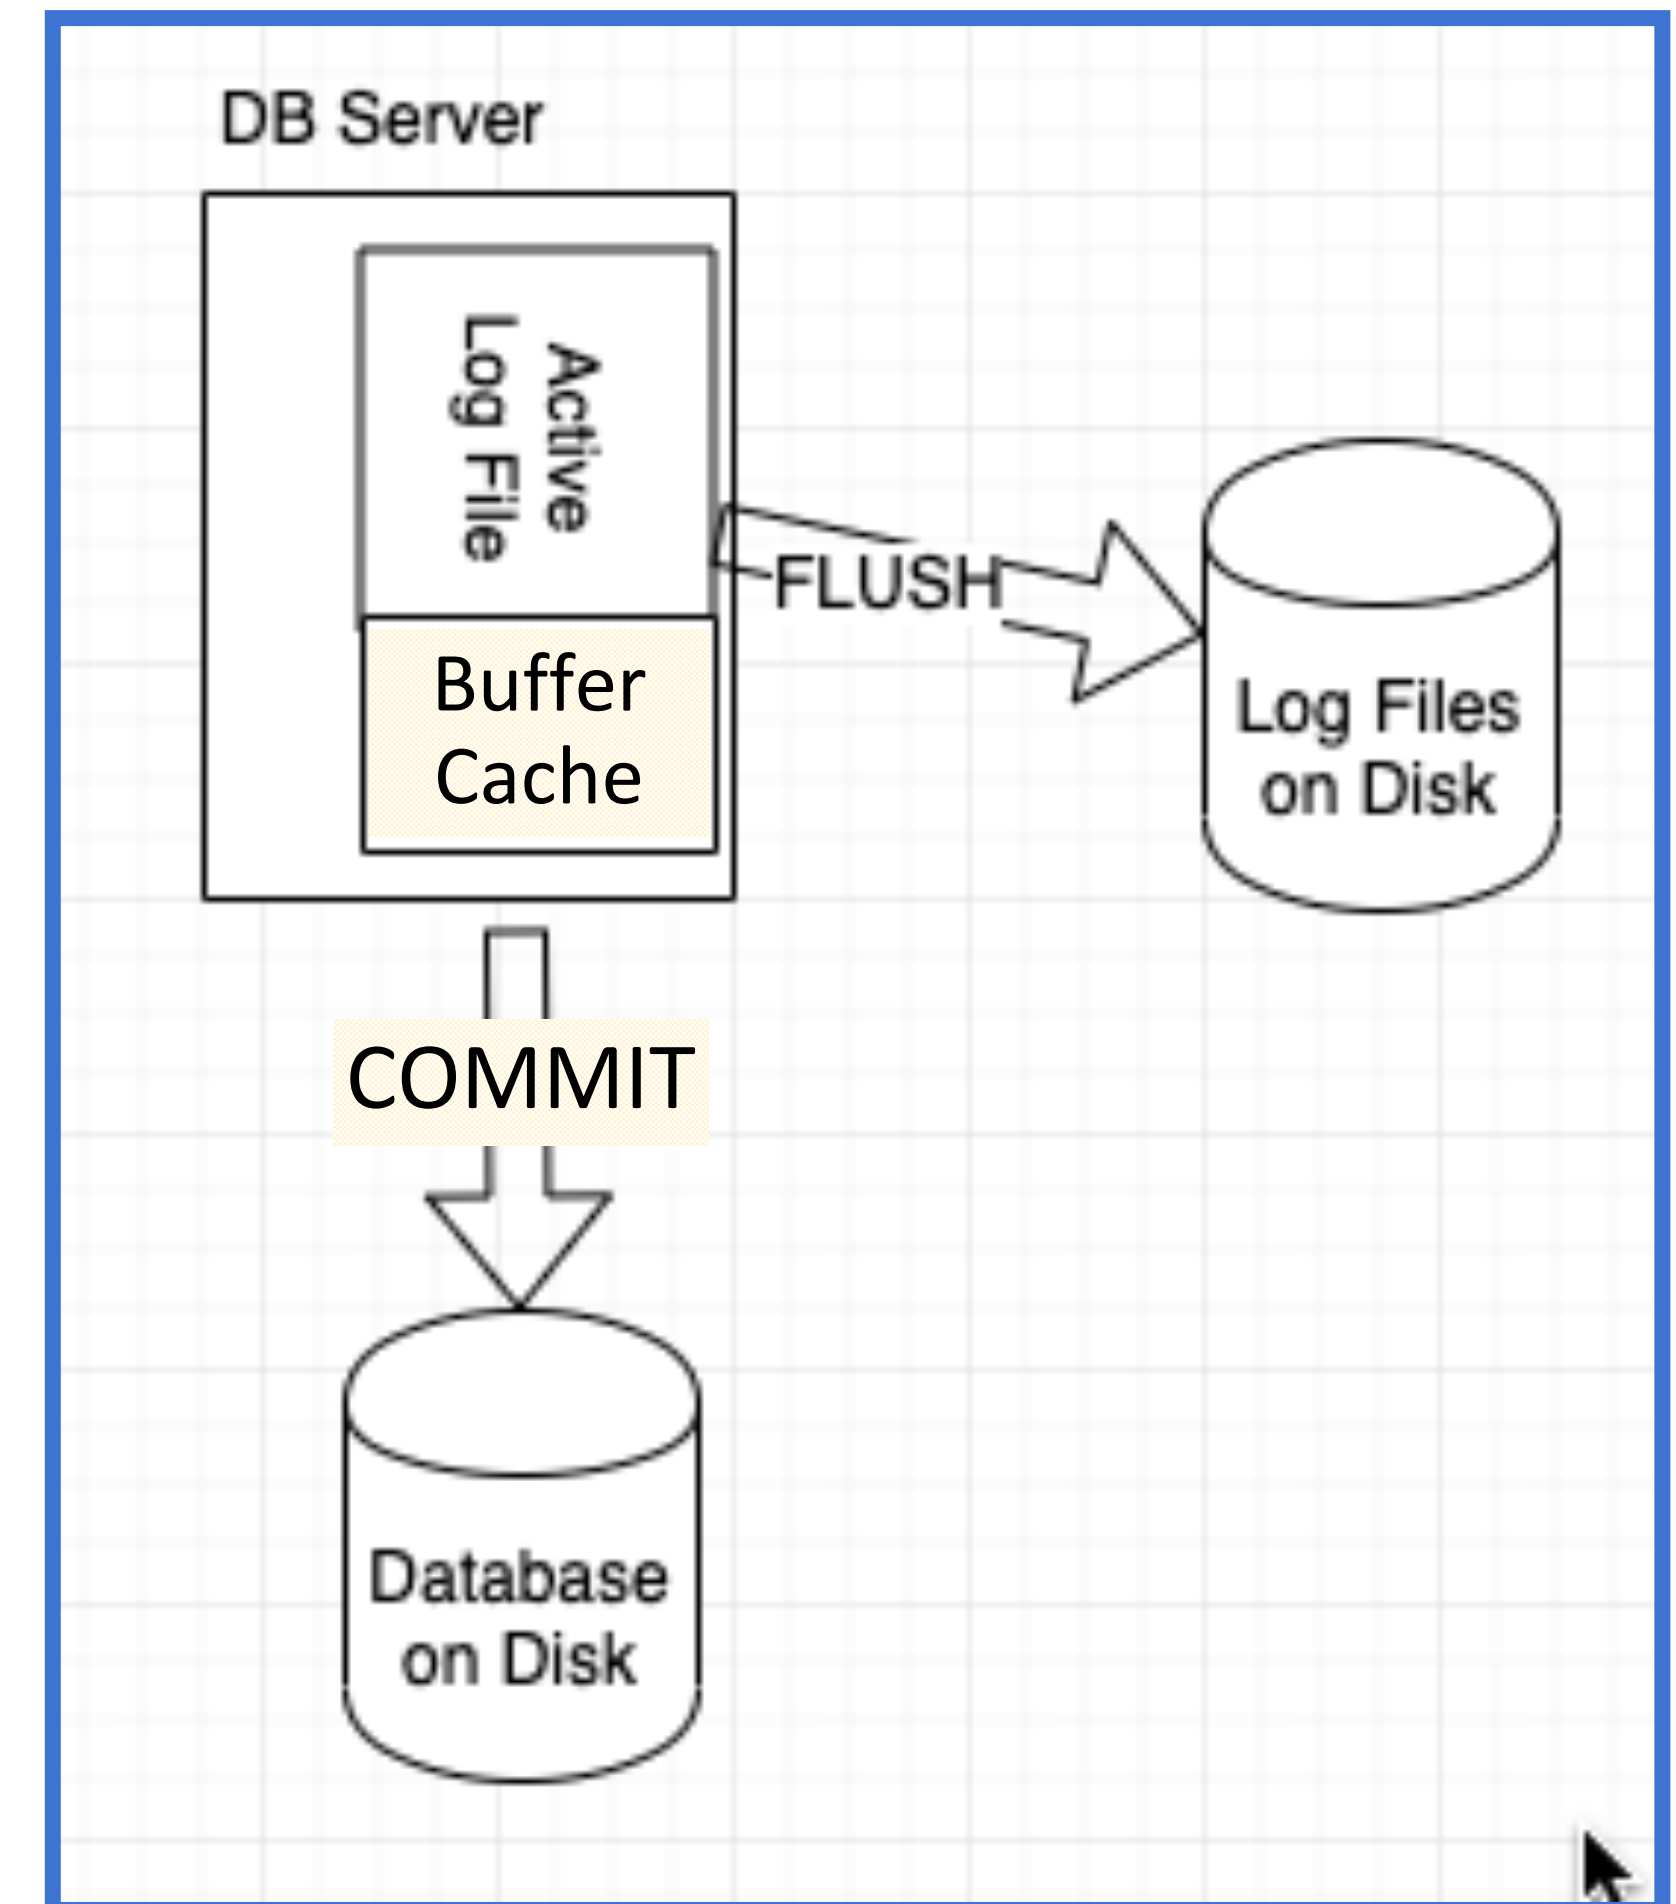
**Deadlock:** Occurs when two transactions are each waiting on a resource that the other transaction holds.

# Concurrency Control

## How can this happen?

All database updates happen within cache, and are only written out to the database on disk when a COMMIT is issued.

Transactions can access uncommitted data that resides in cache, if allowed.



# Concurrency Control

## Sequence of events without conflict

1. Read Item 500
2. Update Item 500
3. Write Item 500

User A

Processing order within the DB Server

1. Read Item 500 for A (into Cache)
2. Read Item 800 for B (into Cache)
3. Update Item 500 for A (in Cache)
4. Update Item 800 for B (in Cache)
5. Write Item 500 for A (commit)
6. Write Item 800 for B (commit)

1. Read Item 800
2. Update Item 800
3. Write Item 800

User B

# Concurrency Control

## Sequence of events with conflict – Lost Update

1. Read Item 500  
("count" = 120)
2. Set "count" up by 12
3. Write Item 500

User A

NOTE:

A's change in Step 3 is LOST

Processing order within the DB Server

1. Read Item 500 for A (into Cache)
2. Read Item 500 for B (into Cache)
3. Update "count" to 132 for A  
(in Cache)
4. Write Item 500 for A (commit)
5. Update "count" to 110 for B  
(in Cache)
6. Write Item 500 for B (commit),  
overlayed "count" 132 with 110

1. Read Item 500  
("count" = 120)
2. Set "count" down by 10
3. Write Item 500

User B

# Concurrency Control

**How does the DBMS Software handle concurrency problems?**

**Resource locking** prevents multiple transactions from obtaining copies of the same record when the record is about to be changed.



# Concurrency Control

**Implicit locks** are locks automatically placed by the DBMS

**Explicit locks** are issued by an application program issuing SQL "lock" commands (rare)

(Explicit locks issued by application code can increase the risk of deadlocks.)



# Concurrency Control

**Lock granularity** refers to size of the locked resource:

- Row (most granular)
- Page ...
- Table ...
- Database (least granular)

**Row level lock granularity minimizes the amount of data being locked, and so improves overall concurrency and processing throughput.**

**All locks require storage and processing time, so less granular locks can reduce locking overhead.**

# Concurrency Control

## Types of locks:

- An **exclusive lock** prohibits other users from reading or updating the locked resource
- A **shared lock** allows other users to read the locked resource, but they cannot update it.

# Concurrency Control

## Serializable Transactions

Refers to two transactions that execute in parallel, but the result is the same as if the transactions had run separately, leaving the database in a consistent state.

The transactions are **fully isolated** from each other.

# Concurrency Control

## Serialized Transactions

1. Lock Item 500
2. Read Item 500  
("count" = 120)
3. Set "count" up by 12
4. Write Item 500
5. Release lock on Item 500

User A

1. Lock Item 500
2. Read Item 500
3. Set "count" down by 10
4. Write Item 500
5. Release lock on Item 500

User B

# Concurrency Control

## Serialized Transactions

Processing order within the DB Server

1. Lock Item 500 for A
2. Read Item 500 for A
3. Lock Item 500 for B – CANNOT grab lock, Trans B goes into WAIT STATE until the lock on Item 500 is released
4. Update "count" to 132 (for A)
5. Write Item 500 for A (commit)
6. Release lock on Item 500 from A
7. Lock Item 500 for B
8. Read Item 500 for B
9. Update "count" to 122 (for B)
10. Write Item 500 for B (commit)
11. Release lock on Item 500 from B

# Concurrency Control

**Deadlock**, or the "deadly embrace", occurs when two transactions are each waiting on a resource that the other transaction holds.

## **Breaking deadlock:**

- Almost every DBMS has algorithms for detecting deadlock
- When a deadlock occurs, the DBMS aborts one of the transactions (selected randomly), allowing one transaction to complete, while the aborted transaction rolls back any partially completed work.

**The aborted transaction must be resubmitted by the application**

# Concurrency Control

## Deadlock

1. Lock Invoice 2022-1
2. Read Invoice 2022-1
3. Lock Item 500
4. Read Item 500
5. Update Invoice 2022-1 (commit)
6. Update Item 500 (commit)
7. Release lock on Invoice 2022-1
8. Release lock on Item 500

User A

1. Lock Item 500
2. Read Item 500
3. Lock Invoice 2022-1
4. Read Invoice 2022-1
5. Update Item 500 (commit)
6. Update Invoice 2022-1 (commit)
7. Release lock on Item 500
8. Release lock on Invoice 2022-1

User B



# Concurrency Control

## Deadlock

Processing order within the DB Server

1. Lock Invoice 2022-1 for A
2. Read Invoice 2022-1 for A
3. Lock Item 500 for B
4. Read Item 500 for B
4. Lock Item 500 for A – Cannot grab lock, Trans A goes into WAIT state
5. Lock Invoice 2022-1 for B - Cannot grab lock, Trans B goes into WAIT state
6. Deadlock is detected
7. DBMS selects ONE transaction (Trans B arbitrarily) to abort
8. Release lock on Item 500 from B due to abort
9. Lock Item 500 for A
10. Update Invoice 2022-1, Item 500 for A (commit)
11. Release locks on Invoice 2022-1 and Item 500 from Trans A

# Concurrency Control

The acronym **ACID**

DBMS Transaction management must be **A**tomic, **C**onsistent, **I**solated, and **D**urable.

**Atomic** means either *ALL* database updates within a transaction are committed or *ALL* updates are rolled back.

**Consistent** means the DBMS software will not allow any updates to violate any database constraints. Once committed, the transaction leaves the data in a consistent state.

# Concurrency Control

**Isolated** means the DBMS software was designed and written in such a way that database administrators can declare the desired isolation level to control concurrency.

**Durable** means that once a transaction is committed to disk, it remains committed, even in the case of a system failure.

# Concurrency Control

The ANSI SQL-92 standard defines four levels of **transaction isolation**:

- **Read uncommitted** – allows reads of uncommitted changes (i.e. “dirty” reads, unrepeatable reads and phantom reads)
- **Read committed** – allows unrepeatable reads and phantom reads, but no dirty reads
- **Repeatable read** – allows phantom reads, but no unrepeatable reads
- **Serializable** – guarantees no read anomalies

# Concurrency Control

## Anomalies:

**“Dirty Read”** – a READ request is allowed to read from cache the updated but uncommitted copy of the data

**“Unrepeatable Read”** – One READ request following another READ request gets two different results.

# Concurrency Control

## Anomalies:

**“Phantom Read”** – a READ request returns a different set of rows at different times. For example the first SELECT returns 10 rows; the next SELECT returns 11 rows (including a “phantom” row) that was inserted after the first read.

# Concurrency Control

Why should we care about **Isolation** levels?

## Isolation levels impact

- The speed at which processing get done
- The frequency and duration of various locks

## Trade-offs

- Faster throughput versus data consistency and reliability




# Concurrency Control

Why do we care about **ACID**?


- Relational DBMS software is written with extremely complex algorithms to ensure ACID compliance
- Ensuring ACID compliance has **trade-offs**
  - Do I want the fastest possible processing? That is, applications/users never have to wait for commits, locks, etc.
  - Do I want to be sure that my application/users maintains full data integrity – that is, no risk of inconsistent data?

# Concurrency Control

The DBMS software itself determines where you land on this scale.



NoSQL databases



Relational Databases

# Concurrency Control

## Summary

- As multiple transactions run against the database concurrently, anomalies can arise
  - Lost Updates, Inconsistent Reads, Deadlocks
- The DBMS software uses locks to minimize anomalies arising from concurrent updates
- Relational DBMS software was written to adhere to ACID transaction compliance