

Module1

April 16, 2022

0.0.1 Grading

The final score that you will receive for your programming assignment is generated in relation to the total points set in your programming assignment item—not the total point value in the nbgrader notebook. When calculating the final score shown to learners, the programming assignment takes the percentage of earned points vs. the total points provided by nbgrader and returns a score matching the equivalent percentage of the point value for the programming assignment. **DO NOT CHANGE VARIABLE OR METHOD SIGNATURES** The autograder will not work properly if you change the variable or method signatures.

0.0.2 Validate Button

Please note that this assignment uses nbgrader to facilitate grading. You will see a **validate button** at the top of your Jupyter notebook. If you hit this button, it will run tests cases for the lab that aren't hidden. It is good to use the validate button before submitting the lab. Do know that the labs in the course contain hidden test cases. The validate button will not let you know whether these test cases pass. After submitting your lab, you can see more information about these hidden test cases in the Grader Output. *Cells with longer execution times will cause the validate button to time out and freeze. Please know that if you run into Validate time-outs, it will not affect the final submission grading.*

1 Part 1. Data cleaning and Exploratory Data Analysis (EDA)

This part will practice data cleaning and Exploratory Data Analysis (EDA) using a house price dataset and mpg dataset. The first dataset is from a Kaggle competition (<https://www.kaggle.com/c/house-prices-advanced-regression-techniques/overview>), where the task is to predict a house sale price given house features.

[282]:

```
!ls
```

data Module1.ipynb

[283]:

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import math
```

1.1 1. Import data and visually inspect the table [9 pts]

1.1.1 1a) Data import and basic inspection. [5 pts]

We can import the csv data using `pd.read_csv()` function. We can use `df.head()` and `df.tail()` to show the first and last 5 entries. `df.iloc[[3,5,7]]` shows the entries corresponding to the index 3,5,7. What is the maximum value of the feature `MSSubClass` among the last 10 entries? Update the value of `maxval` to the correct integer value.

```
[284]: df = pd.read_csv('data/house_data.csv') #it is the same data as the kaggle
      ↪ competition's train.csv.
      # your code here

      # uncomment maxval and update the correct integer value
      maxval = max(df['MSSubClass'].tail(10))
      maxval
```

[284]: 180

```
[285]: # this cell tests that you correctly updated maxval
```

1.1.2 1b) `df.info()` gives the overview of the data frame. Inspect the data using `df.info()` and answer below questions. [4 pts]

1b-i) Which column is the target?

1b-ii) How many features are in the data? Exclude the target. (Id is not a useful feature, but let's still include)

1b-iii) How many observations (samples) are in the data?

1b-iv) How many features have null values based on the data overview?

```
[286]: ## Fixed 1b2

      # your code here
      df.info()

      # uncomment and update to the correct string value
      # copy directly from the uneditd df column name (e.g., 'LandContour')
      ANS_1b1 = 'SalePrice'
      # uncomment and update to the correct integer value
      ANS_1b2 = 80
      # uncomment and update to the correct integer value
      ANS_1b3 = 1460
```

```
# uncomment and update to the correct integer value
ANS_1b4 = sum(df.isna().any())
ANS_1b4
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1460 non-null   int64
1   MSSubClass            1460 non-null   int64
2   MSZoning              1460 non-null   object
3   LotFrontage          1201 non-null   float64
4   LotArea               1460 non-null   int64
5   Street               1460 non-null   object
6   Alley                91 non-null     object
7   LotShape              1460 non-null   object
8   LandContour           1460 non-null   object
9   Utilities             1460 non-null   object
10  LotConfig             1460 non-null   object
11  LandSlope             1460 non-null   object
12  Neighborhood          1460 non-null   object
13  Condition1            1460 non-null   object
14  Condition2            1460 non-null   object
15  BldgType              1460 non-null   object
16  HouseStyle            1460 non-null   object
17  OverallQual           1460 non-null   int64
18  OverallCond           1460 non-null   int64
19  YearBuilt             1460 non-null   int64
20  YearRemodAdd          1460 non-null   int64
21  RoofStyle             1460 non-null   object
22  RoofMatl              1460 non-null   object
23  Exterior1st           1460 non-null   object
24  Exterior2nd           1460 non-null   object
25  MasVnrType            1452 non-null   object
26  MasVnrArea            1452 non-null   float64
27  ExterQual             1460 non-null   object
28  ExterCond             1460 non-null   object
29  Foundation            1460 non-null   object
30  BsmtQual              1423 non-null   object
31  BsmtCond              1423 non-null   object
32  BsmtExposure          1422 non-null   object
33  BsmtFinType1          1423 non-null   object
34  BsmtFinSF1            1460 non-null   int64
35  BsmtFinType2          1422 non-null   object
36  BsmtFinSF2            1460 non-null   int64
37  BsmtUnfSF             1460 non-null   int64
```

38	TotalBsmtSF	1460 non-null	int64
39	Heating	1460 non-null	object
40	HeatingQC	1460 non-null	object
41	CentralAir	1460 non-null	object
42	Electrical	1459 non-null	object
43	1stFlrSF	1460 non-null	int64
44	2ndFlrSF	1460 non-null	int64
45	LowQualFinSF	1460 non-null	int64
46	GrLivArea	1460 non-null	int64
47	BsmtFullBath	1460 non-null	int64
48	BsmtHalfBath	1460 non-null	int64
49	FullBath	1460 non-null	int64
50	HalfBath	1460 non-null	int64
51	BedroomAbvGr	1460 non-null	int64
52	KitchenAbvGr	1460 non-null	int64
53	KitchenQual	1460 non-null	object
54	TotRmsAbvGrd	1460 non-null	int64
55	Functional	1460 non-null	object
56	Fireplaces	1460 non-null	int64
57	FireplaceQu	770 non-null	object
58	GarageType	1379 non-null	object
59	GarageYrBlt	1379 non-null	float64
60	GarageFinish	1379 non-null	object
61	GarageCars	1460 non-null	int64
62	GarageArea	1460 non-null	int64
63	GarageQual	1379 non-null	object
64	GarageCond	1379 non-null	object
65	PavedDrive	1460 non-null	object
66	WoodDeckSF	1460 non-null	int64
67	OpenPorchSF	1460 non-null	int64
68	EnclosedPorch	1460 non-null	int64
69	3SsnPorch	1460 non-null	int64
70	ScreenPorch	1460 non-null	int64
71	PoolArea	1460 non-null	int64
72	PoolQC	7 non-null	object
73	Fence	281 non-null	object
74	MiscFeature	54 non-null	object
75	MiscVal	1460 non-null	int64
76	MoSold	1460 non-null	int64
77	YrSold	1460 non-null	int64
78	SaleType	1460 non-null	object
79	SaleCondition	1460 non-null	object
80	SalePrice	1460 non-null	int64

dtypes: float64(3), int64(35), object(43)

memory usage: 924.0+ KB

[286]: 19

```
[287]: # this cell will test your solutions to the four questions above
```

1.2 2. Inspect Null values [16 pts]

The empty values in the data are called null values. Null values can take different forms. Have a look at below example. `np.nan` and `None` are native null values in python. They get displayed differently in the pandas dataframe (`pd.DataFrame`) though. But there are other data types such as empty list, empty dictionary, etc and string values that literally says “null” or that are empty spaces. Depending on how messy the data is, sometimes the table may have null values of one or more kinds, and those can be cleaned manually or automatically if you can write a code to include all possible cases which meanings are null values.

```
[288]: a = [np.nan, None, [], {}, 'NaN', 'Null', 'NULL', 'None', 'NA', '?', '-', '.', '', ' ',  
          '→', ' ', ' ']  
nulldemo = pd.DataFrame(a)  
nulldemo
```

```
[288]:      0  
0    NaN  
1    None  
2     []  
3     {}  
4    NaN  
5    Null  
6    NULL  
7    None  
8     NA  
9      ?  
10     -  
11     .  
12  
13  
14
```

`.isnull()` method applied to pandas dataframe or series can detect null values. `.dropna()` method in pandas will detect null values and can be specified to drop either rows or columns that contain null values. Below shows that `.isnull()` only detects the python-native null values and cannot detect other forms (string value) of variables that meant null.

```
[289]: nulldemo.isnull()
```

```
[289]:      0  
0    True  
1    True  
2   False  
3   False
```

```
4 False
5 False
6 False
7 False
8 False
9 False
10 False
11 False
12 False
13 False
14 False
```

Also, sometimes the python-native null values can have an odd data type such as numpy float.

```
[290]: print(df['MasVnrArea'].iloc[234], df['MasVnrArea'].iloc[234].dtype,
        ↪type(df['MasVnrArea'].iloc[234]))
print(df['MasVnrArea'].isnull().iloc[234])
print(np.isnan(df['MasVnrArea'].iloc[234]))
print(math.isnan(df['MasVnrArea'].iloc[234]))
print(df['MasVnrArea'].iloc[234]==np.nan)
print(df['MasVnrArea'].iloc[234]==np.float64(np.nan))
```

```
nan float64 <class 'numpy.float64'>
True
True
True
False
False
```

`np.isnan()` and `math.isnan()` can detect the nan values with numpy float type, but they will cause errors with native `None` or a string value. Uncomment one of below (one at a time) and run. You'll see error messages.

```
[291]: # your code here

# print(np.isnan(None))
# print(np.isnan('None'))
# print(math.isnan(None))
# print(math.isnan('None'))
```

1.2.1 2a) Check null values type [5 pts]

Let's check if our data has clean null values (one kind) or messy null values (multiple different representations). Run the codes below and visually inspect the printed results. Which column has string-typed null/none values and how many elements are string-typed null/none values?

```
[292]: # prints number of null values detected by .isnull() and string none
for c in df.columns:
    string_null = np.array([x in a[2:] for x in df[c]])
    print(c, df[c].isnull().sum(), string_null.sum())
```

```
Id 0 0
MSSubClass 0 0
MSZoning 0 0
LotFrontage 259 0
LotArea 0 0
Street 0 0
Alley 1369 0
LotShape 0 0
LandContour 0 0
Utilities 0 0
LotConfig 0 0
LandSlope 0 0
Neighborhood 0 0
Condition1 0 0
Condition2 0 0
BldgType 0 0
HouseStyle 0 0
OverallQual 0 0
OverallCond 0 0
YearBuilt 0 0
YearRemodAdd 0 0
RoofStyle 0 0
RoofMatl 0 0
Exterior1st 0 0
Exterior2nd 0 0
MasVnrType 8 864
MasVnrArea 8 0
ExterQual 0 0
ExterCond 0 0
Foundation 0 0
BsmtQual 37 0
BsmtCond 37 0
BsmtExposure 38 0
BsmtFinType1 37 0
BsmtFinSF1 0 0
BsmtFinType2 38 0
BsmtFinSF2 0 0
BsmtUnfSF 0 0
TotalBsmtSF 0 0
Heating 0 0
HeatingQC 0 0
CentralAir 0 0
```

```

Electrical 1 0
1stFlrSF 0 0
2ndFlrSF 0 0
LowQualFinSF 0 0
GrLivArea 0 0
BsmtFullBath 0 0
BsmtHalfBath 0 0
FullBath 0 0
HalfBath 0 0
BedroomAbvGr 0 0
KitchenAbvGr 0 0
KitchenQual 0 0
TotRmsAbvGrd 0 0
Functional 0 0
Fireplaces 0 0
FireplaceQu 690 0
GarageType 81 0
GarageYrBlt 81 0
GarageFinish 81 0
GarageCars 0 0
GarageArea 0 0
GarageQual 81 0
GarageCond 81 0
PavedDrive 0 0
WoodDeckSF 0 0
OpenPorchSF 0 0
EnclosedPorch 0 0
3SsnPorch 0 0
ScreenPorch 0 0
PoolArea 0 0
PoolQC 1453 0
Fence 1179 0
MiscFeature 1406 0
MiscVal 0 0
MoSold 0 0
YrSold 0 0
SaleType 0 0
SaleCondition 0 0
SalePrice 0 0

```

Which column has string-typed null/none values?

```

[293]: # your code here

# uncomment and update to the correct string value
col = 'MasVnrType'

```

How many elements are string-typed null/none values?


```
[294]: # your code here

# uncomment and update to the correct string value
string_null_count = 864
```

```
[295]: # this cell will test your answer about the column with string-typed null/none
      ↪ values
# and the number of string-typed null/none values
```

1.2.2 2b) Inspect observations (rows) with null values. How many observations have at least one missing value? [5 pts]

```
[296]: # your code here

# uncomment and update to the correct integer value
rows_with_nulls = df.isnull().any(axis=1).sum()
rows_with_nulls
```

```
[296]: 1460
```

```
[297]: # this cell will test your answer about the number of rows with null values
```

1.2.3 2c) Make a histogram of null counts [6 pts]

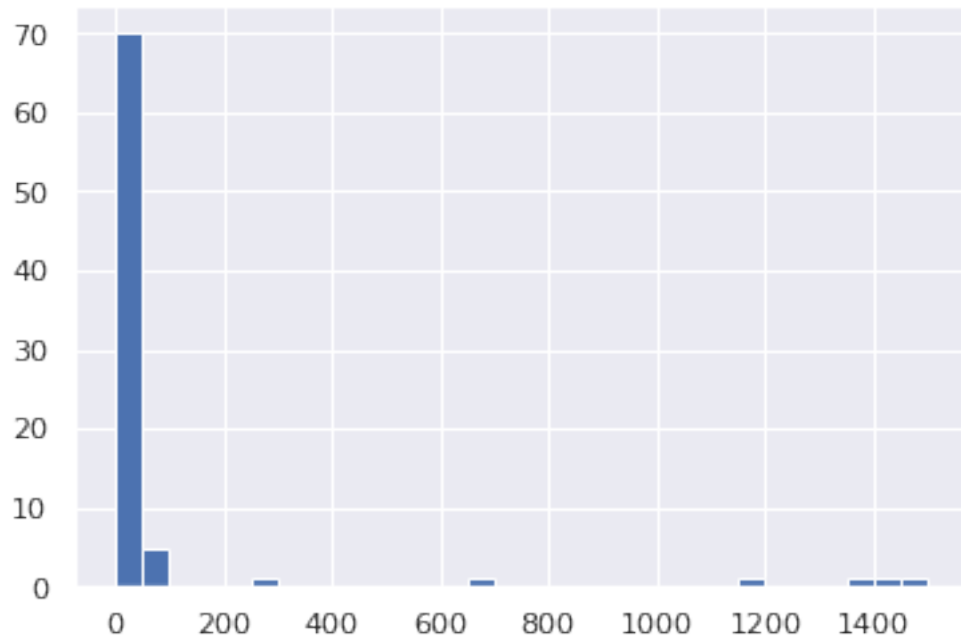
```
[298]: # your code here

# Please uncomment and update
# do not change the names of the variables from null_counts and histogram
null_counts=pd.Series([])
histogram = None # replace the histogram to be the plt.hist() object.

null_counts = df.isnull().sum(axis=0)
histogram = plt.hist(null_counts, bins=np.arange(0,1550, 50))
histogram
# Hint: Use .isnull() and sum over True values on columns.
# You can make it as short as 2-3 lines of code
```

```
[298]: (array([70.,  5.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
          1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.,  0.,
          0.,  1.,  1.,  1.]),
array([  0,   50,  100,  150,  200,  250,  300,  350,  400,  450,  500,
        550,  600,  650,  700,  750,  800,  850,  900,  950, 1000, 1050,
        1100, 1150, 1200, 1250, 1300, 1350, 1400, 1450, 1500])),
```

<a list of 30 Patch objects>)



```
[299]: type(null_counts)
```

```
[299]: pandas.core.series.Series
```

```
[300]: # hidden test 1; tests null_counts
```

```
[301]: # hidden test 2; tests histogram
```

1.3 3. Imputing missing values [33 pts]

In this part, we will decide methods to clean the data with missing values.

Complete case analysis (CCA) is to drop any observations (rows) that have null values. It is suitable if the number of observations with null values are very small (say, less than 5%) compared to the total number of observations.

If the data has a large number of features (columns) and the model(s) does not need that many features (some models work better with less number of features), we can consider dropping features that have many missing values. Before dropping features, it is generally a good idea checking whether the feature with missing values is important feature or not (which may need the analyst's judgement). If the feature is very important for the prediction task (for example, a house size when predicting house price) but has a large amount of missing values, we cannot simply drop the feature, or in a rare case, it could mean that the data is not suitable for the analysis. One will have to work with only the observations that has values on that feature given the number of observations

is sufficient, or collect more data. If we know that those features are not very important and have a large number of missing values, we can drop the features. As a rule of thumb, features with missing values more than either 5% or 10% can be dropped.

1.3.1 3a) Is the data suitable for complete case analysis or not? [5 pts]

```
[302]: # your code here

# uncomment and update to string 'no' or 'yes'
suitable_cca = 'no'
```

```
[303]: # tests solution for whether data is suitable for CCA
```

1.3.2 3b) Dropping feature columns [20 pts]

Let's assume we want to keep columns that have null values 5% or less and discard any column that has null values more than 5%. Treat the string type "None" as a category and not null value. ##### 3b-i) According to above condition, how many features can be kept and imputed? [5 pts] ##### 3b-ii) Which columns have null values 5% or less of total, so we can impute? [5 pts] ##### 3b-iii) Which columns have null values more than 5% of total, so we should throw? [5 pts]

```
[304]: # your code here
perc_missing = df.isnull().sum() / len(df)
perc_df = pd.DataFrame({
    'Column_Name' : df.columns,
    'Perc_Missing' : perc_missing
})
# Complete the codes below by uncommenting and changing the values of
# → features_to_impute and features_to_throw.
# Each should be a list of feature names (e.g. ['LotFrontage', 'Alley', ...]). Do
# → not change the variable names.
# There are hidden tests which will grade above three questions.
perc_df.sort_values('Perc_Missing', inplace=True, ascending=False)
perc_df.head(20)

features_to_throw = list(perc_df['Column_Name'][:11])
features_to_impute = perc_df['Column_Name'][11:19]
print(len(features_to_impute), features_to_impute)
print(len(features_to_throw), features_to_throw)
```

```
8 BsmtFinType2    BsmtFinType2
BsmtExposure     BsmtExposure
BsmtQual         BsmtQual
BsmtCond         BsmtCond
BsmtFinType1     BsmtFinType1
```

```

MasVnrArea      MasVnrArea
MasVnrType      MasVnrType
Electrical      Electrical
Name: Column_Name, dtype: object
11 ['PoolQC', 'MiscFeature', 'Alley', 'Fence', 'FireplaceQu', 'LotFrontage',
    'GarageYrBlt', 'GarageCond', 'GarageType', 'GarageFinish', 'GarageQual']

```

```
[305]: # Hidden test for 3b-i
```

```
[306]: # Hidden test for 3b-ii
```

```
[307]: # Hidden test for 3b-iii
```

1.3.3 3b-iv) Remove the columns according to the above result. Replace the df with the new result. Also remove Id column as it's not a useful feature. [5 pts]

```

[308]: # your code here

# remove the columns according to the above result, replace df with the new
↳ results
# also remove ID column as it's not a useful feature
features_to_throw.append('Id')
df.drop(features_to_throw, inplace=True, axis=1)
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 69 columns):
#   Column                Non-Null Count  Dtype
---  -
0   MSSubClass             1460 non-null  int64
1   MSZoning               1460 non-null  object
2   LotArea                1460 non-null  int64
3   Street                 1460 non-null  object
4   LotShape               1460 non-null  object
5   LandContour            1460 non-null  object
6   Utilities              1460 non-null  object
7   LotConfig              1460 non-null  object
8   LandSlope              1460 non-null  object
9   Neighborhood           1460 non-null  object
10  Condition1             1460 non-null  object
11  Condition2             1460 non-null  object
12  BldgType               1460 non-null  object
13  HouseStyle             1460 non-null  object
14  OverallQual            1460 non-null  int64
15  OverallCond            1460 non-null  int64

```

16	YearBuilt	1460	non-null	int64
17	YearRemodAdd	1460	non-null	int64
18	RoofStyle	1460	non-null	object
19	RoofMatl	1460	non-null	object
20	Exterior1st	1460	non-null	object
21	Exterior2nd	1460	non-null	object
22	MasVnrType	1452	non-null	object
23	MasVnrArea	1452	non-null	float64
24	ExterQual	1460	non-null	object
25	ExterCond	1460	non-null	object
26	Foundation	1460	non-null	object
27	BsmtQual	1423	non-null	object
28	BsmtCond	1423	non-null	object
29	BsmtExposure	1422	non-null	object
30	BsmtFinType1	1423	non-null	object
31	BsmtFinSF1	1460	non-null	int64
32	BsmtFinType2	1422	non-null	object
33	BsmtFinSF2	1460	non-null	int64
34	BsmtUnfSF	1460	non-null	int64
35	TotalBsmtSF	1460	non-null	int64
36	Heating	1460	non-null	object
37	HeatingQC	1460	non-null	object
38	CentralAir	1460	non-null	object
39	Electrical	1459	non-null	object
40	1stFlrSF	1460	non-null	int64
41	2ndFlrSF	1460	non-null	int64
42	LowQualFinSF	1460	non-null	int64
43	GrLivArea	1460	non-null	int64
44	BsmtFullBath	1460	non-null	int64
45	BsmtHalfBath	1460	non-null	int64
46	FullBath	1460	non-null	int64
47	HalfBath	1460	non-null	int64
48	BedroomAbvGr	1460	non-null	int64
49	KitchenAbvGr	1460	non-null	int64
50	KitchenQual	1460	non-null	object
51	TotRmsAbvGrd	1460	non-null	int64
52	Functional	1460	non-null	object
53	Fireplaces	1460	non-null	int64
54	GarageCars	1460	non-null	int64
55	GarageArea	1460	non-null	int64
56	PavedDrive	1460	non-null	object
57	WoodDeckSF	1460	non-null	int64
58	OpenPorchSF	1460	non-null	int64
59	EnclosedPorch	1460	non-null	int64
60	3SsnPorch	1460	non-null	int64
61	ScreenPorch	1460	non-null	int64
62	PoolArea	1460	non-null	int64
63	MiscVal	1460	non-null	int64

```

64  MoSold          1460 non-null   int64
65  YrSold          1460 non-null   int64
66  SaleType        1460 non-null   object
67  SaleCondition    1460 non-null   object
68  SalePrice        1460 non-null   int64
dtypes: float64(1), int64(34), object(34)
memory usage: 787.2+ KB

```

```
[309]: # tests that you properly updated df
```

1.3.4 3c) Impute missing data [8 pts]

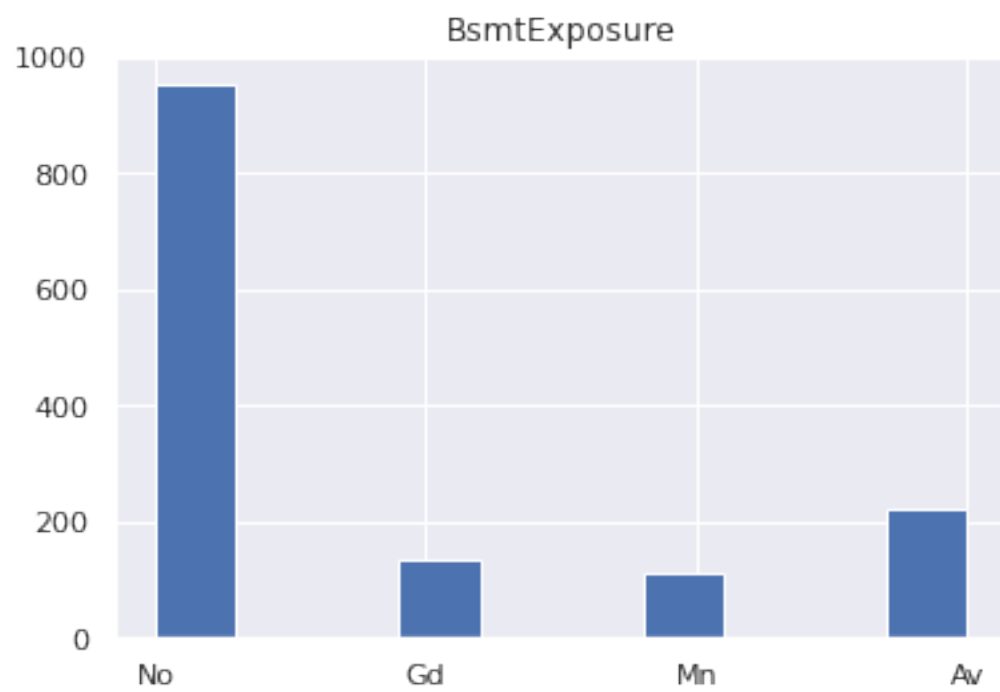
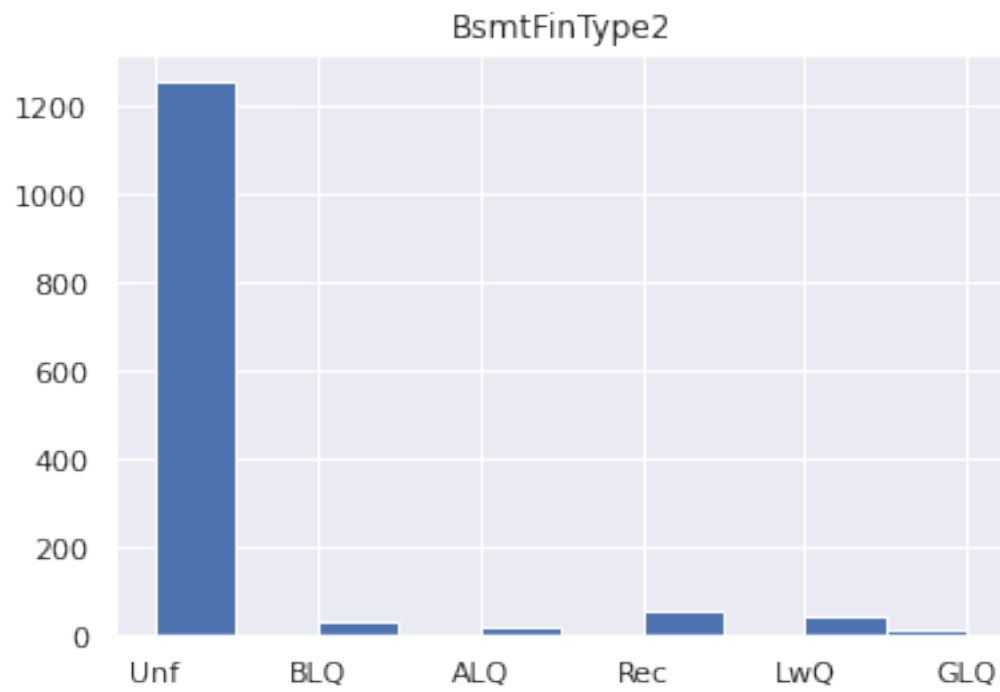
Before imputing columns, we need to think about what methods to use to impute columns. The imputation strategy can be different depending on the variable types and variable value distribution. There are many imputation techniques, but let's use a few simple ones.

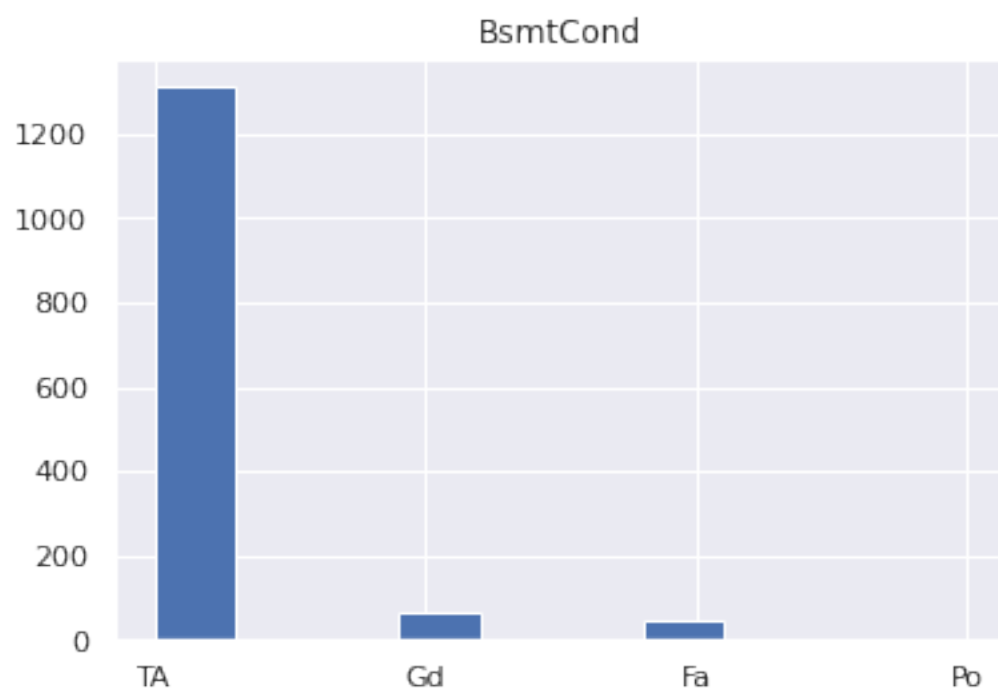
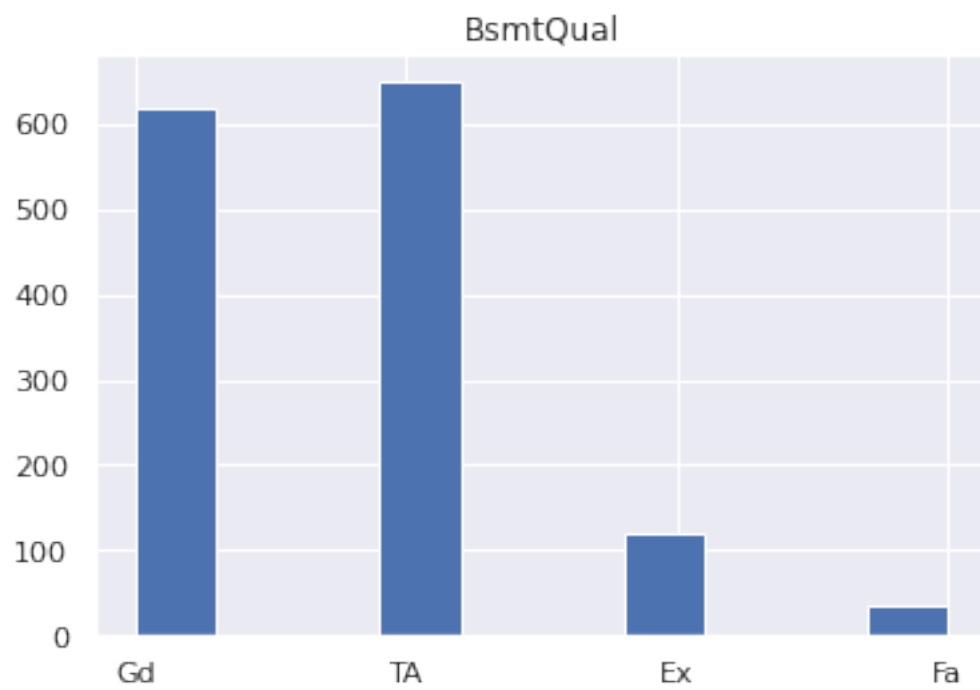
For a numerical variable imputation, we impute mean value if the distribution is symmetric while we use median value to impute when the distribution is skewed. Another method is to assign an arbitrary value that's outside the normal range. Though it can be useful to capture missingness, but it can create outliers. Both mean/median and arbitrary imputation methods are simple to use and suitable when missing values are 5% (no more than 10%) as a rule of thumb. Both methods can distort the original distribution.

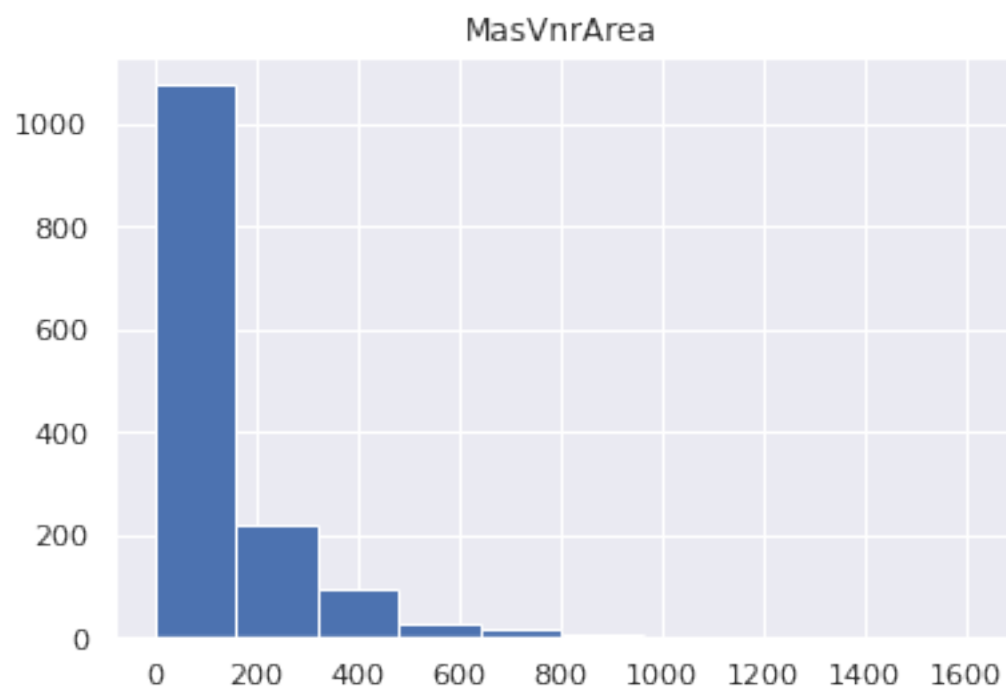
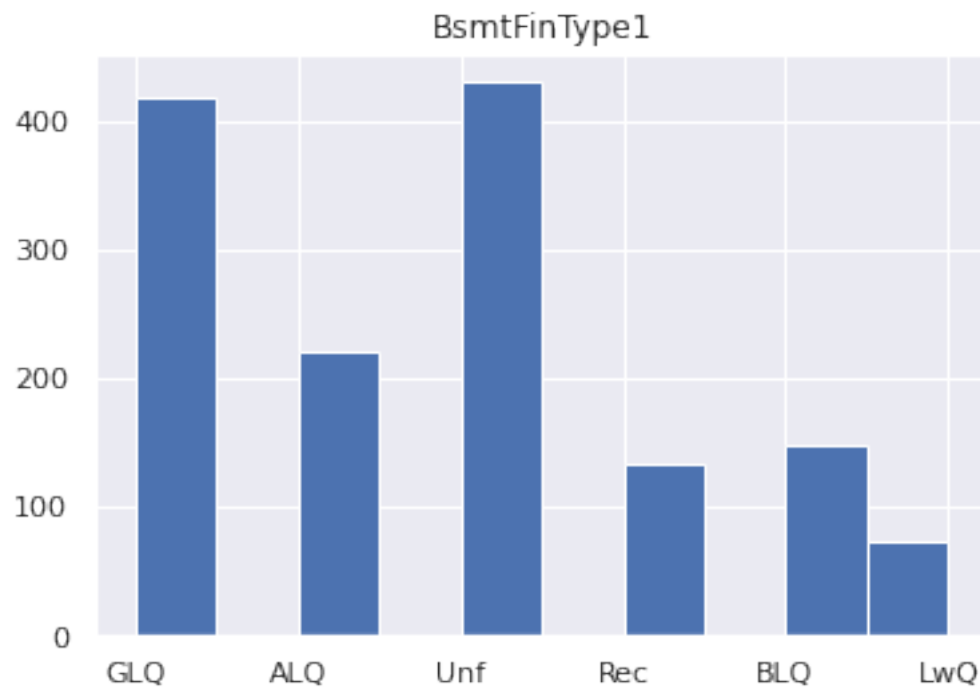
For a categorical variable imputation, we can impute with the most frequent categorical value. It is a simple method but it can distort the original distribution. It is also possible to create a "missing" category to capture missingness. The advantage of using missing category is that it captures missingness but its disadvantage is that it creates another rare category.

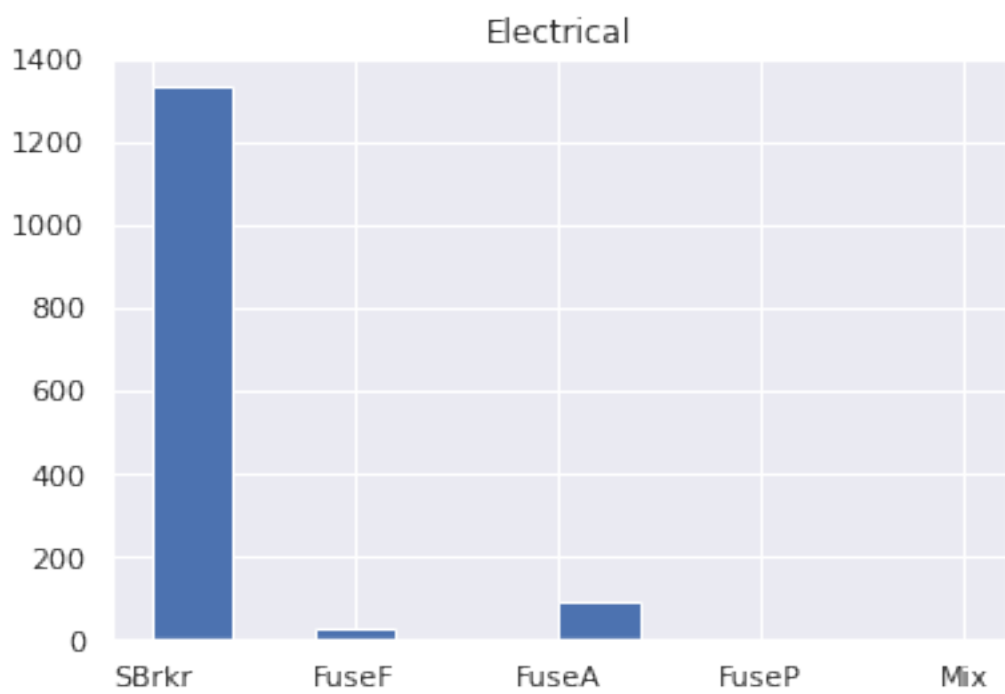
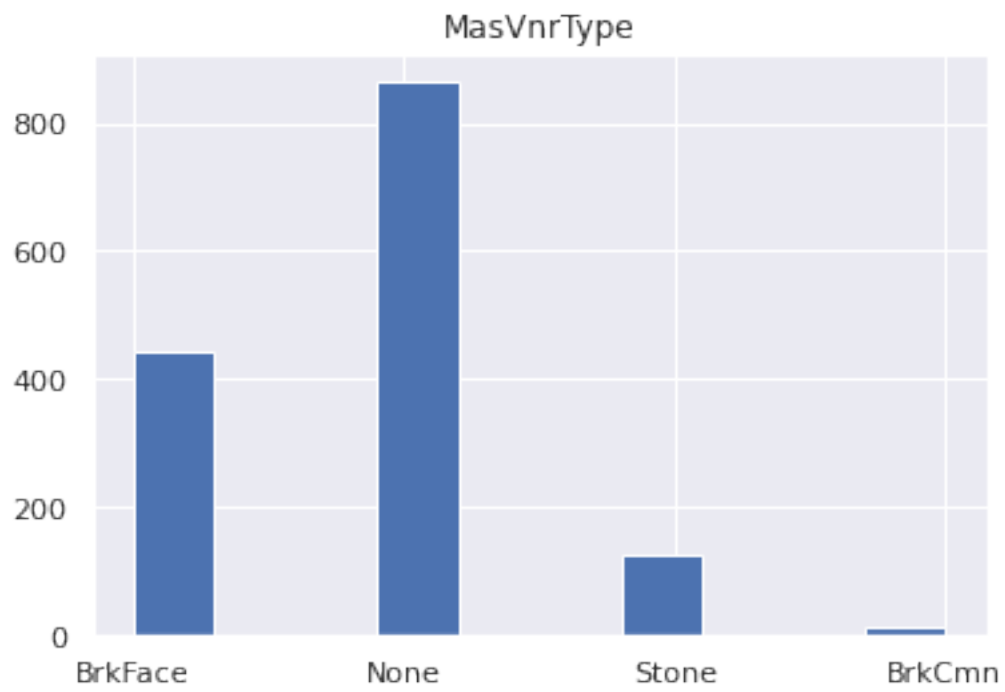
Below code shows histograms of feature columns that we can impute.

```
[310]: for c in features_to_impute:
        df[c].hist()
        plt.title(c)
        plt.show()
```









1.3.5 3c-i) Impute missing data for features in features_to_impute. Choose an appropriate method among mean or median imputation methods for numerical variable(s) and frequentest value imputation for categorical variable(s). [8 pts]

You can inspect variable types by eyes, or use below code as a help. Replace those columns with imputed values. Do not change the column name or the data frame name. Do not add new columns to the data frame.

Hint: You can use `.mode()` function to find the most frequent value in a Series.

Hint: You may use `.fillna()` function on each feature Series.

```
[311]: for c in features_to_impute:
        print(c, len(df[c].unique()), df[c].dtype)
```

```
BsmtFinType2 7 object
BsmtExposure 5 object
BsmtQual 5 object
BsmtCond 5 object
BsmtFinType1 7 object
MasVnrArea 328 float64
MasVnrType 5 object
Electrical 6 object
```

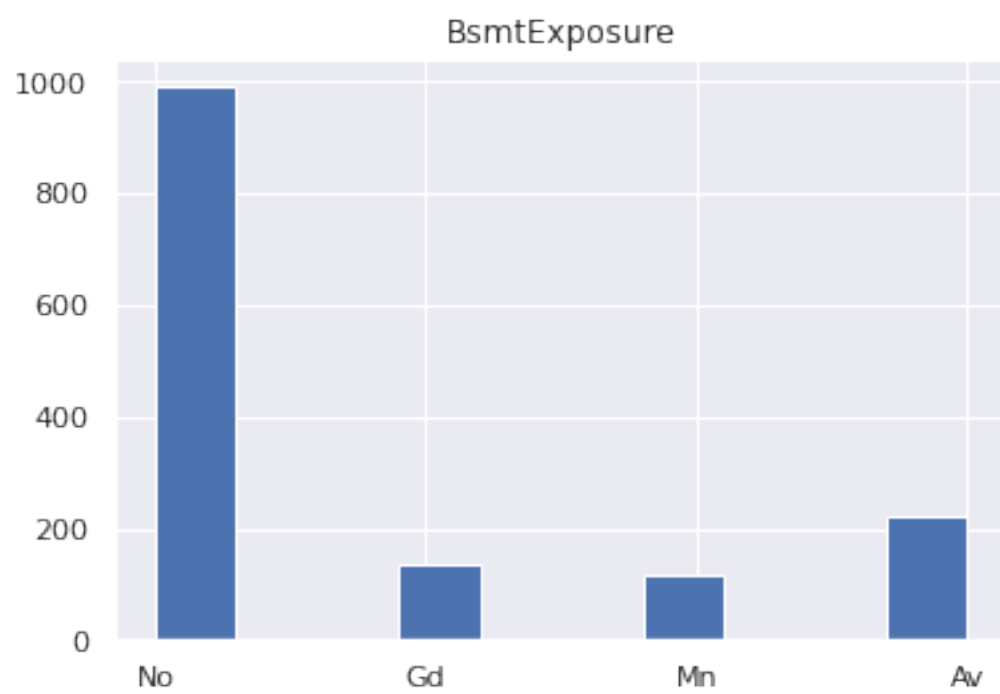
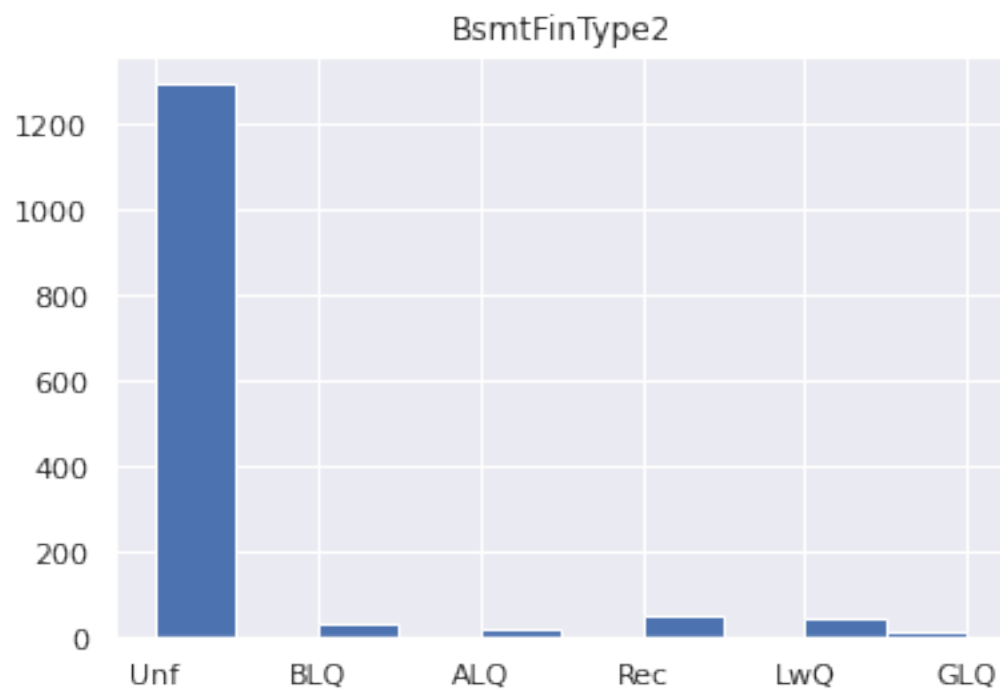
```
[312]: # your code here
from statistics import mode
print(df['BsmtFinType2'].dtype)
# use this cell for potential debugging
```

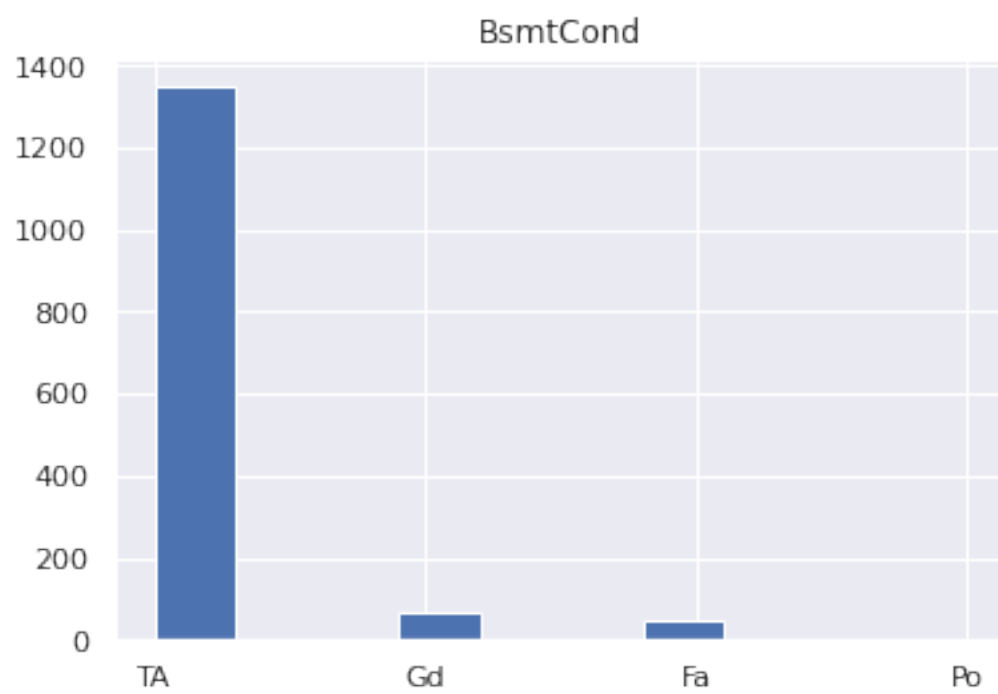
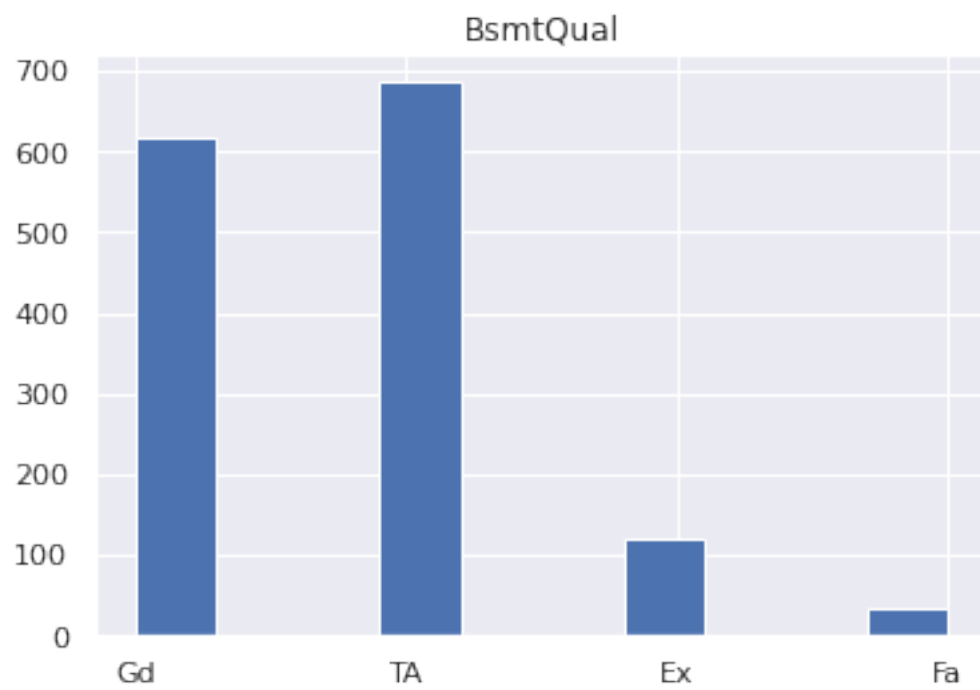
```
object
```

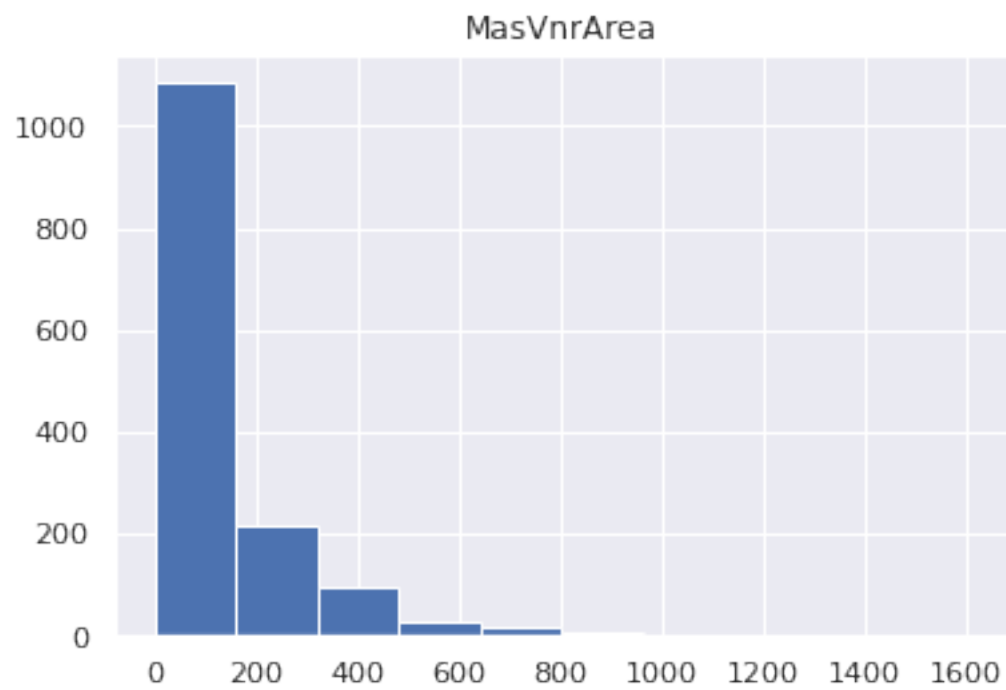
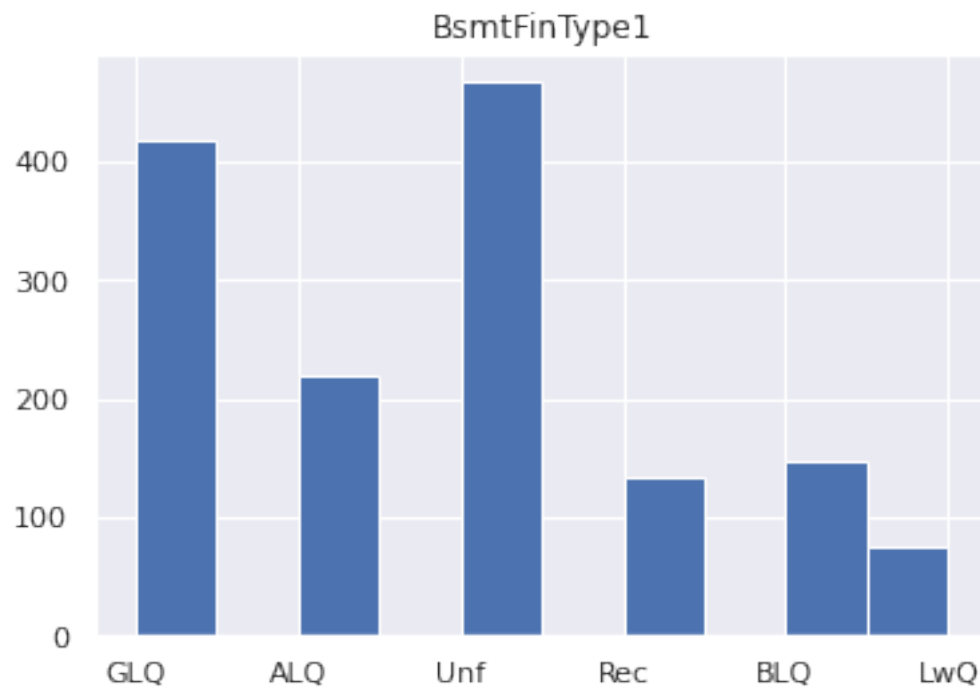
```
[313]: # impute missing data
# your code here

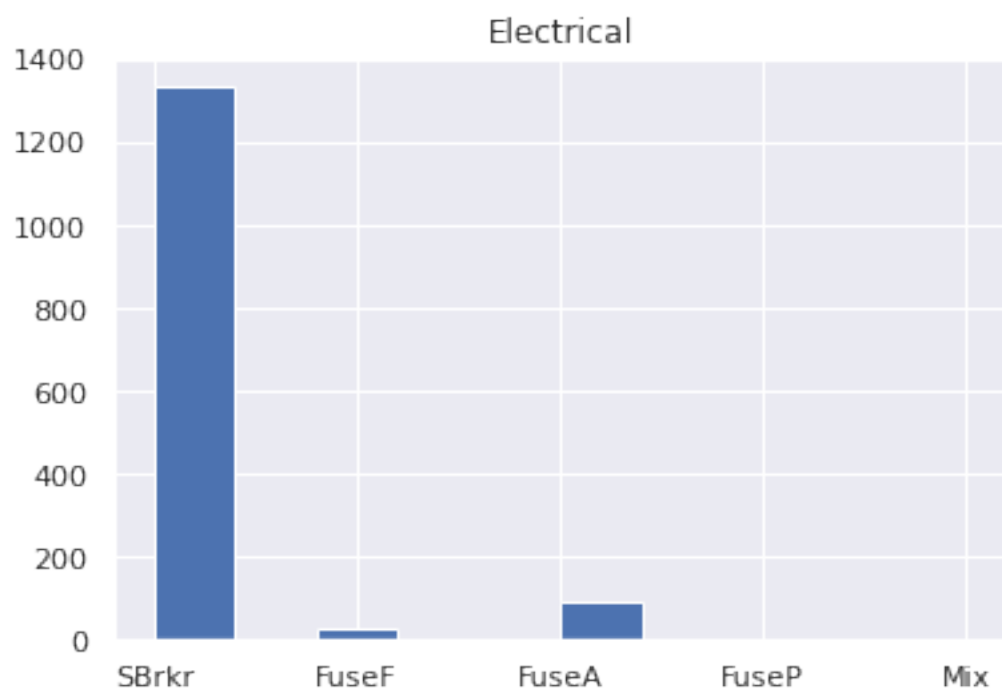
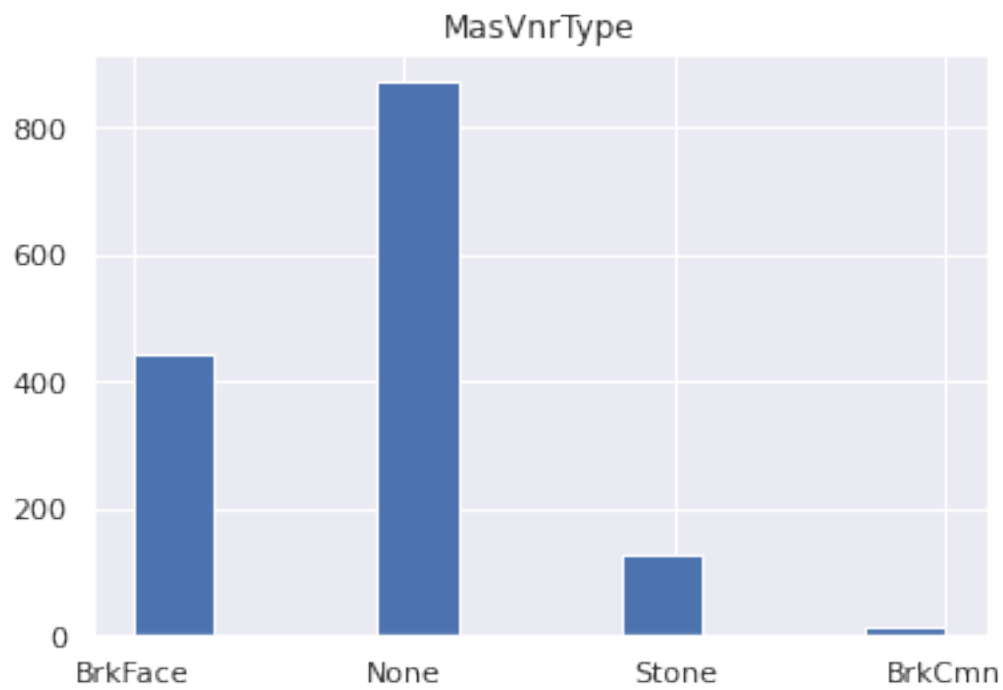
for col in features_to_impute:
    if df[col].dtype == 'object':
        df[col] = df[col].fillna(mode(df[col]))
    if df[col].dtype == 'float64':
        df[col] = df[col].fillna(df[col].median())
```

```
[314]: for c in features_to_impute:
        df[c].hist()
        plt.title(c)
        plt.show()
```









```
[315]: # tests 'MasVnrType' and 'MasVnrArea'
```

```
[316]: # tsts 'BsmtQual' and 'BsmtCond'
```

```
[317]: # tests 'BsmtExposure' and 'BsmtFinType1'
```

```
[318]: # tests 'BsmtFinType2' and 'Electrical'
```

2 Part 2. EDA, Simple Linear Regression

In this part, we will use a simplified data and create a simple linear regression model. The dataset can be downloaded from <https://www.kaggle.com/harlfoxem/housesalesprediction/download>. This dataset contains house sale prices for Kings County, which includes Seattle. It includes homes sold between May 2014 and May 2015. There are several versions of the data. Some additional information about the columns is available here: <https://geodacenter.github.io/data-and-lab/KingCounty-HouseSales2015/>, some of which are copied below.

Variable	Description
id	Identification
date	Date sold
price	Sale price
bedrooms	Number of bedrooms
bathrooms	Number of bathrooms
sqft_liv	Size of living area in square feet
sqft_lot	Size of the lot in square feet
floors	Number of floors
waterfront	'1' if the property has a waterfront, '0' if not.
view	An index from 0 to 4 of how good the view of the property was
condition	Condition of the house, ranked from 1 to 5
grade	Classification by construction quality which refers to the types of materials used and the quality of workmanship. Buildings of better quality (higher grade) cost more to build per unit of measure and command higher value.
sqft_above	Square feet above ground
sqft_basmt	Square feet below ground
yr_built	Year built
yr_renov	Year renovated. '0' if never renovated
zipcode	5 digit zip code
lat	Latitude
long	Longitude
sqft_liv15	Average size of interior housing living space for the closest 15 houses, in square feet
sqft_lot15	Average size of land lost for the closest 15 houses, in square feet


```
[319]: import scipy as sp
import scipy.stats as stats
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import copy
# Set color map to have light blue background
sns.set()
import statsmodels.formula.api as smf
import statsmodels.api as sm
%matplotlib inline
```

```
[320]: df2 = pd.read_csv('data/house_data_washington.csv')
```

2.1 4. Munging data [15 pts]

2.1.1 4a) Date string to numbers [5 pts]

Inspect the data frame and data type of each column. The column 'date' is the date sold, and has string value. We will extract year and month information from the string. In the data frame df2, create new features 'sales_year' and 'sales_month'.

```
[321]: # extract year and month info from the string
# create new features 'sales_year' and 'sales_month' in df2
df2['sales_year'] = df2.date.apply(lambda x: int(x[:4]))
df2['sales_month'] = df2.date.apply(lambda x: int(x[4:6]))
print(df2.groupby('sales_month')['id'].count())
print(df2.groupby('sales_year')['id'].count())
```

```
sales_month
1      978
2     1250
3     1875
4     2231
5     2414
6     2180
7     2211
8     1940
9     1774
10    1878
11    1411
12    1471
Name: id, dtype: int64
sales_year
2014    14633
```

2015 6980
Name: id, dtype: int64

Which month has the most number of sales?

```
[322]: # your code here

# uncomment and update string with capitalized month, e.g., 'December'
most_sales = 'May'
```

Which months has the least number of sales?

```
[323]: # your code here

# uncomment and update string with capitalized month, e.g., 'December'
least_sales = 'January'
```

```
[324]: # tests solutions for most_sales and least_sales
```

```
[325]: df2.head(10)
```

```
[325]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
0	7129300520	20141013T000000	221900	3	1.00	1180	
1	6414100192	20141209T000000	538000	3	2.25	2570	
2	5631500400	20150225T000000	180000	2	1.00	770	
3	2487200875	20141209T000000	604000	4	3.00	1960	
4	1954400510	20150218T000000	510000	3	2.00	1680	
5	7237550310	20140512T000000	1225000	4	4.50	5420	
6	1321400060	20140627T000000	257500	3	2.25	1715	
7	2008000270	20150115T000000	291850	3	1.50	1060	
8	2414600126	20150415T000000	229500	3	1.00	1780	
9	3793500160	20150312T000000	323000	3	2.50	1890	

	sqft_lot	floors	waterfront	view	...	sqft_basement	yr_built	\
0	5650	1.0	0	0	...	0	1955	
1	7242	2.0	0	0	...	400	1951	
2	10000	1.0	0	0	...	0	1933	
3	5000	1.0	0	0	...	910	1965	
4	8080	1.0	0	0	...	0	1987	
5	101930	1.0	0	0	...	1530	2001	
6	6819	2.0	0	0	...	0	1995	
7	9711	1.0	0	0	...	0	1963	
8	7470	1.0	0	0	...	730	1960	
9	6560	2.0	0	0	...	0	2003	

	yr_renovated	zipcode	lat	long	sqft_living15	sqft_lot15	\
0	0	98178	47.5112	-122.257	1340	5650	
1	1991	98125	47.7210	-122.319	1690	7639	
2	0	98028	47.7379	-122.233	2720	8062	

3	0	98136	47.5208	-122.393	1360	5000
4	0	98074	47.6168	-122.045	1800	7503
5	0	98053	47.6561	-122.005	4760	101930
6	0	98003	47.3097	-122.327	2238	6819
7	0	98198	47.4095	-122.315	1650	9711
8	0	98146	47.5123	-122.337	1780	8113
9	0	98038	47.3684	-122.031	2390	7570

	sales_year	sales_month
0	2014	10
1	2014	12
2	2015	2
3	2014	12
4	2015	2
5	2014	5
6	2014	6
7	2015	1
8	2015	4
9	2015	3

[10 rows x 23 columns]

```
[326]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21613 non-null  int64
1   date                 21613 non-null  object
2   price               21613 non-null  int64
3   bedrooms            21613 non-null  int64
4   bathrooms           21613 non-null  float64
5   sqft_living         21613 non-null  int64
6   sqft_lot            21613 non-null  int64
7   floors              21613 non-null  float64
8   waterfront          21613 non-null  int64
9   view                21613 non-null  int64
10  condition            21613 non-null  int64
11  grade               21613 non-null  int64
12  sqft_above          21613 non-null  int64
13  sqft_basement       21613 non-null  int64
14  yr_built            21613 non-null  int64
15  yr_renovated        21613 non-null  int64
16  zipcode             21613 non-null  int64
17  lat                 21613 non-null  float64
```

```

18 long          21613 non-null float64
19 sqft_living15 21613 non-null int64
20 sqft_lot15    21613 non-null int64
21 sales_year    21613 non-null int64
22 sales_month   21613 non-null int64
dtypes: float64(4), int64(18), object(1)
memory usage: 3.8+ MB

```

2.1.2 4b) Variable types [5 pts]

Inspect each feature's data type and variable type. What is the best description for the variable type of following features? Update the string to 'numeric' or 'categorical'.

```

[327]: # your code here

# uncomment the feaures below and update the strings with 'numeric' or
↳ 'categorical'
price = 'numeric'
bathrooms = 'numeric'
waterfront = 'categorical'
grade = 'numeric'
zipcode = 'categorical'
sales_year = 'numeric'

```

```

[328]: # tests that you selected correct variable type for the features in 4b

```

```

[329]: # this part is ungraded, but useful to run to check
# your code here

for c in df2.columns[2:]:
    print(c, df2[c].unique())

```

```

price [ 221900  538000  180000 ...  610685 1007500  402101]
bedrooms [ 3  2  4  5  1  6  7  0  8  9 11 10 33]
bathrooms [1.  2.25 3.  2.  4.5  1.5  2.5  1.75 2.75 3.25 4.  3.5  0.75 4.75
 5.  4.25 3.75 0.  1.25 5.25 6.  0.5  5.5  6.75 5.75 8.  7.5  7.75
 6.25 6.5 ]
sqft_living [1180 2570  770 ... 3087 3118 1425]
sqft_lot [ 5650  7242 10000 ...  5813  2388  1076]
floors [1.  2.  1.5 3.  2.5 3.5]
waterfront [0 1]
view [0 3 4 2 1]
condition [3 5 4 1 2]
grade [ 7  6  8 11  9  5 10 12  4  3 13  1]
sqft_above [1180 2170  770 1050 1680 3890 1715 1060 1890 1860  860 1430 1370
1810

```

1980	1600	1200	1250	2330	2270	1070	2450	1710	1750	1400	790	2570	2320
1190	1510	1090	1280	930	2360	890	2620	2600	3595	1570	920	3160	990
2290	2165	1640	1000	2130	2830	2250	2420	3250	1850	1590	1260	2519	1540
1110	1770	2720	2240	3070	2380	2390	880	1040	910	3450	2350	1900	1010
960	2660	1610	765	3520	1290	1960	1160	1210	1270	1440	2190	2920	1460
1170	1240	3140	2030	2310	700	1080	2520	2780	1560	1450	1720	2910	1620
1360	2070	2460	1390	2140	1320	1340	1550	940	1380	3670	2370	1130	980
3540	2500	1760	1030	1780	3400	2680	1670	2590	820	1220	2440	2090	1100
1330	1420	1690	2150	1910	1350	1940	900	1630	2714	850	1870	1950	2760
2020	1120	1480	1230	2280	3760	3530	830	1300	2740	1830	720	2010	3360
800	1730	760	1700	4750	5310	580	2653	2850	2210	2630	3500	1740	1140
2160	2650	970	2040	2180	2220	1660	3370	2690	1930	3150	3030	2050	2490
2560	1275	2580	560	1820	1840	2990	3230	1580	3480	2510	1410	2120	3300
3840	1500	1530	2840	833	2000	6070	950	2200	4040	1920	1490	3470	3130
2610	3260	2260	430	3390	630	4860	3860	2810	870	3180	2770	4030	4410
2400	1520	3040	6050	4740	1970	5403	3350	3580	1790	750	2860	2750	2340
2870	4120	3200	2550	1805	4150	1384	2060	2110	3590	2100	2540	1880	1150
1470	1255	1800	4370	3190	2730	4570	2470	670	2900	4670	4230	2156	1020
2940	2640	2710	3100	3610	4270	840	3090	2300	380	2480	3460	3060	3064
3000	1654	2790	1310	2230	2430	3680	2670	2208	810	740	1422	490	2080
3440	5670	4475	730	3410	3010	600	2960	3570	4300	3990	780	3020	5990
440	4460	4190	2800	2530	1650	3690	2932	3720	4250	3110	2963	4930	2950
5000	2452	2820	1981	640	2495	2403	5320	6720	660	2341	4210	3830	3280
2980	5153	1990	1646	610	710	5450	3504	3210	1782	2930	590	4280	680
3880	3430	3750	4130	5710	3380	3330	4700	3220	3362	3510	3810	620	4490
2410	3050	1008	3488	4070	3420	5770	1605	520	1088	3555	4360	3960	2700
4340	1552	3850	2303	3270	4350	3640	2174	4160	2496	5180	5130	6350	3770
2153	3780	2890	1714	2201	2970	992	3950	3527	2835	3915	1427	4870	3340
3620	4310	3930	4080	5400	570	3310	6110	3320	3490	3859	3710	1798	4600
3560	3940	3600	3800	1105	2305	3290	5050	1556	1553	4000	1657	3001	4220
480	3120	3740	530	3700	5230	5370	3080	4140	4430	3550	1159	1288	2880
4610	1122	3052	1479	7680	3820	1934	5080	2675	2506	5760	2154	4390	3240
1995	1689	2782	2395	4400	6200	3526	4320	2483	4380	4580	4180	2064	3650
1726	2019	4240	1256	500	1355	1747	1678	1833	1414	4115	3597	3170	390
1976	5830	2601	3920	2641	5070	2518	3910	3660	3695	4020	2803	2074	2038
4060	4890	2329	1264	1095	690	4090	1392	2844	902	4560	2811	4720	2168
5610	2683	4900	2095	4290	4050	4260	4440	6220	1175	998	2356	4500	3900
3831	1315	4470	4810	2286	2927	4760	8570	5140	1679	1811	2849	1676	1757
3730	2441	2163	5250	2795	2415	3970	4200	1068	5240	1509	1954	4820	1651
4100	1752	3630	2885	3154	1129	2632	1996	4010	550	410	6430	3790	2031
1652	2434	3316	1899	2331	2497	2216	4170	1341	1961	5584	8860	2507	5220
4850	5844	5530	2145	650	1982	4910	3605	1778	1463	2783	1946	1358	3870
1864	1845	6290	3980	2382	2979	3674	2726	5440	1295	2115	6085	3265	3136
6640	4620	3361	2245	2242	1078	2577	1329	420	4330	1975	7420	1788	2299
1092	4225	1087	1904	470	2966	2192	2253	5550	4133	4285	1216	540	9410
2075	5330	2166	1628	1808	1352	2557	6380	7880	2734	1363	1769	2093	1677
2588	5190	2298	1491	2961	5020	5980	4540	844	6120	2233	4480	4110	4770
2473	995	5160	1494	2007	1048	3002	4780	2155	2014	4980	2665	4830	4790

```

5010 370 2105 3006 3004 2689 4660 1746 2678 2755 2414 901 4630 2068
2807 2643 2181 4510 4420 1604 1435 3045 2717 2905 4940 5110 2533 6660
3485 2659 5090 2375 1964 866 1595 944 5480 809 5040 1764 1656 1802
460 2692 1544 2044 1212 4083 8020 3905 1502 4590 384 2092 6090 1615
7320 1396 1484 1765 5490 1453 1643 5300 1381 4065 290 1313 5430 1397
2793 2475 1936 3028 798 2575 3276 1584 2393 2029 3222 1072 1785 1984
962 2423 2052 2538 2437 2789 2906 4800 7850 2196 1847 2658 2655 3855
1728 963 2223 1611 2015 2448 1489 1116 3745 1002 3202 1347 1481 2311
2544 2584 2217 3569 3181 1921 2612 2671 2598 3284 3266 1076 2594 2718
1794 2481 3845 1413 1876 3148 2413 1767 5060 806 2547 1834 2024 1165
2134 1741 2798 1852 2099 3216 1094 2891 2432 2283 2701 1658 893 2009
1444 2744 3078 3065 1578 2815 4960 1571 6530 4640 1536 3172 6370 3223
1608 2229 3135 1408 1763 4840 1232 2502 2424 1296 1914 988 3828 3056
2267 1131 2796 1812 1084 2025 1564 1239 2568 1528 2628 2185 2478 2669
1912 2828 2425 1446 3206 2406 1419 2056 1144 2456 4950 3192 828 2529
2732 1987 3906 4073 2578 2738 3691 1061 2846 2542 1889 3336 3236 1451
1983 2313 1824 1322 1766 2301 3274 1108 2864 2716 1572 3281 2656 2398
1867 1613 2587 2623 894 1606 2244 2026 2238 2517 2708 2555 1405 4450
1248 6420 2531 1333 2198 3087 3118 1425]
sqft_basement [ 0 400 910 1530 730 1700 300 970 760 720 700 820 780
790
330 1620 360 588 1510 410 990 600 560 550 1000 1600 500 1040
880 1010 240 265 290 800 540 380 710 840 770 480 570 1490
620 1250 1270 120 650 180 1130 450 1640 1460 1020 1030 750 640
1070 490 1310 630 2000 390 430 850 210 1430 1950 440 220 1160
860 580 2060 1820 1180 200 1150 1200 680 530 1450 1170 1080 960
1100 280 870 460 1400 1320 660 1220 900 420 1580 1380 475 690
270 350 935 1370 980 1470 160 950 50 740 1780 1900 340 470
370 140 1760 130 610 520 890 1110 150 1720 810 190 1290 670
1800 1120 1810 60 1050 940 310 930 1390 1830 1300 510 1330 1590
920 1420 1240 1960 1560 2020 1190 2110 1280 250 2390 1230 170 830
1260 1410 1340 590 1500 1140 260 100 320 1480 1060 1284 1670 1350
2570 2590 1090 110 2500 90 1940 1550 2350 2490 1481 1360 1135 1520
1850 1660 2130 2600 1690 243 1210 2620 1024 1798 1610 1440 1570 1650
704 1910 1630 2360 1852 2090 2400 1790 2150 230 70 1680 2100 3000
1870 1710 2030 875 1540 2850 2170 506 906 145 2040 784 1750 374
518 2720 2730 1840 3480 2160 1920 2330 1860 2050 4820 1913 80 2010
3260 2200 415 1730 652 2196 1930 515 40 2080 2580 1548 1740 235
861 1890 2220 792 2070 4130 2250 2240 894 1990 768 2550 435 1008
2300 2610 666 3500 172 1816 2190 1245 1525 1880 862 946 1281 414
2180 276 1248 602 516 176 225 1275 266 283 65 2310 10 1770
2120 295 207 915 556 417 143 508 2810 20 274 248]
yr_built [1955 1951 1933 1965 1987 2001 1995 1963 1960 2003 1942 1927 1977 1900
1979 1994 1916 1921 1969 1947 1968 1985 1941 1915 1909 1948 2005 1929
1981 1930 1904 1996 2000 1984 2014 1922 1959 1966 1953 1950 2008 1991
1954 1973 1925 1989 1972 1986 1956 2002 1992 1964 1952 1961 2006 1988
1962 1939 1946 1967 1975 1980 1910 1983 1978 1905 1971 2010 1945 1924
1990 1914 1926 2004 1923 2007 1976 1949 1999 1901 1993 1920 1997 1943

```

```

1957 1940 1918 1928 1974 1911 1936 1937 1982 1908 1931 1998 1913 2013
1907 1958 2012 1912 2011 1917 1932 1944 1902 2009 1903 1970 2015 1934
1938 1919 1906 1935]
yr_renovated [ 0 1991 2002 2010 1999 1992 2013 1994 1978 2005 2008 2003 1984
1954
2014 2011 1974 1983 1945 1990 1988 1957 1977 1981 1995 2000 1998 1970
1989 2004 1986 2009 2007 1987 1973 2006 1985 2001 1980 1971 1979 1997
1950 1969 1948 2015 1968 2012 1963 1951 1993 1962 1996 1972 1953 1955
1982 1956 1940 1976 1946 1975 1958 1964 1959 1960 1967 1965 1934 1944]
zipcode [98178 98125 98028 98136 98074 98053 98003 98198 98146 98038 98007 98115
98107 98126 98019 98103 98002 98133 98040 98092 98030 98119 98112 98052
98027 98117 98058 98001 98056 98166 98023 98070 98148 98105 98042 98008
98059 98122 98144 98004 98005 98034 98075 98116 98010 98118 98199 98032
98045 98102 98077 98108 98168 98177 98065 98029 98006 98109 98022 98033
98155 98024 98011 98031 98106 98072 98188 98014 98055 98039]
lat [47.5112 47.721 47.7379 ... 47.3906 47.3339 47.6502]
long [-122.257 -122.319 -122.233 -122.393 -122.045 -122.005 -122.327 -122.315
-122.337 -122.031 -122.145 -122.292 -122.229 -122.394 -122.375 -121.962
-122.343 -122.21 -122.306 -122.341 -122.169 -122.166 -122.172 -122.218
-122.36 -122.314 -122.304 -122.11 -122.07 -122.357 -122.368 -122.157
-122.31 -122.132 -122.362 -122.282 -122.18 -122.027 -122.347 -122.016
-122.364 -122.175 -121.977 -122.371 -122.151 -122.301 -122.451 -122.322
-122.189 -122.384 -122.369 -122.281 -122.29 -122.114 -122.122 -122.116
-122.149 -122.339 -122.335 -122.344 -122.32 -122.297 -122.192 -122.215
-122.16 -122.179 -122.287 -122.036 -122.073 -121.987 -122.125 -122.34
-122.025 -122.008 -122.291 -122.365 -122.199 -122.194 -122.387 -122.372
-122.391 -122.351 -122.386 -122.249 -122.277 -122.378 -121.958 -121.714
-122.08 -122.196 -122.184 -122.133 -122.38 -122.082 -122.109 -122.053
-122.349 -122.295 -122.253 -122.248 -122.303 -122.294 -122.226 -122.266
-122.098 -122.212 -122.244 -122.39 -122.352 -121.85 -122.152 -122.054
-122.072 -121.998 -122.296 -122.299 -122.381 -122.358 -122.128 -122.171
-122.174 -122.026 -122.353 -121.943 -122.286 -122.336 -122.359 -122.162
-122.3 -122.176 -121.996 -122.118 -122.193 -122.023 -122.224 -122.168
-122.231 -122.331 -122.374 -122.182 -122.308 -122.307 -121.999 -122.376
-122.2 -122.039 -122.102 -122.188 -122.379 -122.043 -122.153 -122.191
-122.219 -122.312 -121.911 -121.994 -122.165 -122.37 -122.158 -122.047
-122.284 -122.017 -122.275 -122.268 -122.367 -122.217 -122.373 -122.013
-122.214 -122.034 -122.164 -121.899 -122.183 -121.95 -122.324 -122.216
-122.395 -122.213 -122.345 -122.278 -122.111 -121.711 -122.27 -122.178
-122.147 -121.772 -122.302 -122.438 -122.223 -122.042 -122.323 -122.255
-122.4 -122.261 -122.071 -122.206 -122.272 -122.23 -122.144 -122.143
-122.181 -122.154 -122.311 -122.274 -122.077 -122. -122.298 -122.058
-121.837 -122.333 -122.057 -122.252 -122.093 -122.012 -122.052 -122.354
-122.22 -122.49 -121.875 -122.24 -122.078 -122.173 -121.854 -122.222
-122.28 -122.137 -122.159 -121.974 -122.141 -122.029 -121.709 -122.19
-121.97 -122.329 -122.195 -122.06 -121.959 -122.095 -122.148 -122.146
-122.35 -121.901 -122.241 -122.129 -122.289 -122.305 -122.022 -122.385
-121.779 -122.032 -122.402 -122.482 -122.227 -121.982 -122.161 -122.046

```

-122.156 -122.127 -122.33 -122.197 -122.041 -122.103 -122.318 -122.382
 -122.271 -121.955 -122.211 -122.262 -122.258 -122.121 -122.221 -122.234
 -122.089 -122.123 -122.167 -121.909 -122.107 -122.064 -122.066 -122.062
 -122.264 -122.186 -122.087 -121.88 -121.864 -122.205 -122.363 -122.139
 -122.018 -122.225 -122.285 -122.084 -122.177 -122.056 -122.316 -122.021
 -122.348 -122.009 -122.131 -122.411 -122.198 -122.256 -122.117 -122.097
 -122.075 -121.845 -122.083 -122.259 -121.87 -122.015 -122.007 -121.86
 -122.409 -121.755 -121.972 -122.251 -122.317 -121.776 -122.115 -122.283
 -122.242 -122.001 -122.024 -122.309 -122.113 -121.771 -122.239 -122.273
 -122.396 -122.094 -122.267 -122.326 -122.13 -122.269 -121.853 -122.05
 -122.346 -122.076 -121.826 -122.124 -121.758 -122.202 -121.785 -121.872
 -122.006 -122.004 -122.321 -121.882 -122.101 -122.03 -122.185 -122.1
 -121.759 -121.965 -122.201 -122.366 -122.313 -122.405 -122.02 -122.279
 -122.355 -121.934 -122.15 -122.356 -121.993 -122.044 -122.134 -121.867
 -122.01 -121.991 -122.011 -121.983 -122.228 -122.033 -122.276 -122.119
 -121.937 -122.361 -122.325 -122.203 -122.136 -122.237 -122.209 -122.049
 -122.288 -122.106 -122.037 -122.207 -122.263 -121.915 -122.204 -122.09
 -122.069 -121.852 -121.787 -121.976 -122.377 -122.059 -122.383 -121.989
 -122.019 -122.208 -121.878 -122.328 -122.25 -122.338 -122.388 -122.265
 -122.332 -122.399 -122.397 -122.014 -121.956 -122.092 -122.028 -122.293
 -122.12 -122.035 -122.14 -122.04 -122.112 -121.906 -122.17 -122.238
 -122.512 -121.997 -121.89 -122.463 -121.908 -122.086 -122.389 -121.913
 -122.163 -121.918 -122.108 -122.502 -122.392 -122.236 -121.859 -121.981
 -122.342 -121.96 -121.978 -122.47 -121.91 -121.966 -122.065 -122.246
 -122.41 -121.879 -122.079 -122.099 -122.187 -121.98 -122.002 -122.138
 -121.898 -122.235 -122.126 -121.782 -121.995 -122.401 -121.858 -121.888
 -121.752 -122.063 -122.26 -121.78 -121.708 -121.721 -122.403 -121.945
 -122.243 -122.45 -121.927 -122.085 -122.088 -121.973 -122.055 -122.398
 -121.984 -121.912 -121.903 -121.946 -122.232 -122.412 -122.104 -122.048
 -122.479 -122.155 -121.833 -121.778 -122.003 -121.99 -121.926 -122.051
 -121.986 -122.245 -121.861 -122.431 -121.964 -122.142 -122.074 -122.247
 -122.497 -121.769 -121.827 -121.979 -121.871 -122.091 -121.754 -121.746
 -121.92 -121.992 -122.406 -121.359 -121.789 -121.707 -122.068 -122.404
 -122.334 -121.799 -121.774 -121.985 -121.865 -121.724 -122.415 -121.756
 -121.809 -122.135 -121.691 -122.038 -121.877 -121.94 -121.968 -121.988
 -121.315 -121.902 -122.514 -122.414 -121.883 -121.866 -121.744 -122.096
 -122.061 -121.881 -121.745 -122.461 -122.067 -121.868 -121.646 -121.93
 -122.105 -121.763 -121.718 -121.967 -121.777 -121.957 -121.823 -121.887
 -122.408 -122.462 -122.43 -122.456 -121.897 -121.932 -121.969 -121.916
 -122.081 -121.975 -121.735 -121.801 -121.761 -121.723 -121.924 -122.475
 -121.935 -122.407 -122.448 -122.453 -121.894 -121.936 -121.764 -122.416
 -121.905 -122.464 -121.768 -122.484 -121.738 -121.9 -121.82 -122.455
 -121.889 -122.496 -121.829 -122.505 -121.951 -121.847 -122.509 -121.961
 -121.417 -121.904 -122.503 -121.949 -121.874 -122.432 -121.971 -121.77
 -122.473 -121.896 -121.952 -122.254 -121.743 -121.933 -121.892 -121.749
 -121.473 -121.857 -122.465 -121.838 -121.954 -122.422 -121.931 -121.963
 -122.441 -121.925 -121.352 -122.511 -122.413 -121.876 -121.748 -121.818
 -121.8 -121.929 -121.698 -121.886 -121.802 -121.81 -121.762 -121.781

-121.775 -122.44 -121.773 -121.819 -121.726 -122.459 -122.446 -121.855
 -121.736 -122.499 -122.46 -121.786 -122.421 -121.947 -122.439 -121.834
 -121.804 -122.443 -121.716 -121.848 -122.458 -122.515 -121.922 -121.953
 -121.783 -122.472 -121.944 -121.869 -121.828 -122.452 -121.831 -121.737
 -121.739 -121.863 -121.73 -121.856 -121.747 -121.893 -121.733 -121.846
 -121.821 -121.319 -121.765 -121.75 -122.506 -121.948 -121.921 -122.507
 -122.457 -121.914 -122.469 -121.792 -121.907 -121.841 -121.757 -121.788
 -121.731 -122.449 -121.316 -121.321 -122.504 -121.884 -121.803 -121.842
 -121.719 -121.766 -122.433 -122.519 -121.851 -121.402 -122.454 -122.467
 -121.325 -121.815 -121.676 -121.941 -122.445 -121.76 -121.885 -121.742
 -121.822 -121.895 -121.784 -121.701 -121.713 -121.727 -121.849 -121.835
 -122.435 -122.474 -122.444 -121.939 -121.48 -121.364 -121.767 -122.42
 -121.84 -122.425 -122.447 -121.797 -122.491 -121.917 -121.891 -121.942
 -121.862 -121.725 -121.873 -121.405 -122.486 -121.795 -121.734 -121.403]
 sqft_living15 [1340 1690 2720 1360 1800 4760 2238 1650 1780 2390 2210 1330 1370
 2140
 1890 1610 1060 1280 1400 4110 2240 1220 2200 1030 1760 1860 1520 2630
 2580 1390 1460 1570 2020 1590 2160 1730 1290 2620 2470 2410 3625 1580
 3050 1228 2680 970 1190 1990 1410 1480 2730 1950 2250 2690 2960 2270
 2570 2500 1440 2750 2221 1010 3390 3530 1640 1510 2420 1940 3240 1680
 890 1130 3350 2350 1870 1720 1850 1900 1980 2520 1350 1750 1160 2550
 2370 1240 1270 2990 1380 1540 2090 2640 1830 1620 1880 2340 1710 2700
 3060 2660 1700 1970 1420 2060 2480 1550 1170 2820 1560 2230 2840 1450
 1500 3160 1200 3400 2110 2920 1770 1070 1930 3740 2260 1670 2290 1050
 2540 2190 2030 1230 2330 1300 1430 2770 1250 1630 2590 2130 1100 3836
 1320 2120 3070 1910 2080 1960 2280 1150 3430 2070 2600 830 1260 3120
 2010 1660 1600 2380 3890 4180 2653 2670 3920 2300 2310 2320 3150 1740
 2400 4550 2510 2440 2880 3860 2150 1310 1820 3080 880 2560 3470 1020
 2040 2610 1810 2860 3480 3130 3360 4050 2450 1790 3180 3600 2000 2430
 2850 4680 2360 3930 1490 2460 2077 1920 3630 3220 2100 3230 4300 3850
 2424 2530 3030 2830 2900 2950 1470 940 2740 4210 3340 3980 2180 3715
 2050 1080 2095 1000 3330 2170 1408 1530 2760 3110 950 3000 1307 2220
 4190 3440 3250 1110 2870 1210 2910 1120 4230 1708 3090 3270 2970 1180
 3100 4100 2930 3510 2688 1840 2490 4090 2810 3260 3680 3420 1654 1365
 980 1677 1140 3640 3460 3140 1502 3720 2790 2940 990 2890 860 4750
 1525 3950 5790 760 2234 960 3210 2780 2800 2305 2665 3620 2710 4320
 2650 3370 1509 1277 1981 2434 4640 2242 3040 3970 3200 4600 840 3290
 2214 1162 3010 5600 3820 3540 1975 4800 740 3990 3170 1576 1768 3310
 2980 1429 3900 3380 820 1090 4060 3910 3190 3450 3730 620 3020 3760
 3320 1132 3300 3770 3960 870 3560 4620 3520 1572 3490 1088 3159 4470
 3570 4890 3690 3280 2083 3780 920 1941 1566 850 2496 1040 3410 4240
 4670 4350 1714 5380 4330 3830 5000 2144 1494 1357 930 3580 4250 4080
 3660 1458 3736 1894 2037 1295 4170 3750 3550 4630 1439 3500 2091 900
 3880 3710 1616 720 800 2315 1564 2767 3721 4650 4020 780 1728 2027
 1264 1404 1459 2028 3639 1943 3425 2641 2114 1309 2412 2517 1802 2011
 1466 1414 3193 1845 1156 3670 1696 5340 4440 1745 1884 4690 4920 2406
 4160 3810 4480 2848 1746 2634 2049 5330 1536 2273 3056 4010 4700 910
 2125 1665 2683 3790 700 1855 750 1078 4150 4340 2344 1098 1175 1188

```

3700 3840 4042 2518 3800 2488 3590 2052 810 1528 5030 4740 5070 2967
4280 2724 3610 3940 4940 4770 1811 4830 2876 1805 1216 5170 1304 2474
4590 4130 1492 1364 2168 4140 3543 1303 2005 3650 2583 4310 2451 1448
2955 2142 790 1638 2554 2441 2216 4220 1961 4540 770 4200 3413 1664
2136 3568 4510 1484 1358 2106 1834 2014 4390 4570 2175 6110 4260 710
2112 1934 1518 1302 2622 2619 2382 4290 4560 4000 1336 3112 4070 1468
1571 2605 1138 5110 4850 2165 4410 1678 5610 1984 4660 3870 4370 460
4610 1914 3515 2246 1786 2109 2326 2728 4400 4950 1767 2054 5500 2555
3674 2765 1862 1352 4030 399 2415 2901 1815 2236 2253 2004 1356 2403
1137 1256 4930 4040 2376 4520 4490 2189 2566 2396 1282 2155 1056 2389
2256 3618 1326 1168 4913 806 1369 2405 2875 1425 5220 1442 2333 3335
1321 3045 1546 4730 2697 2822 2076 1757 4780 952 4270 2075 2667 1092
1217 1716 1792 2961 1125 1463 1886 670 4460 2336 3557 5200 2258 1377
2019 2092 4900 2615 1639 1765 1554 1381 4120 5080 1445 2793 2475 998
2384 2575 1398 1584 2439 2197 2029 4362 1443 4420 1691 2495 2437 2547
6210 2009 1847 1346 2578 2879 2255 2815 1608 690 2425 1481 2458 2358
2056 1921 2419 2996 2502 1798 3087 1076 2981 2363 3191 1763 1876 1949
2598 1979 1415 2002 2574 2166 3726 2099 2154 1522 1544 2912 2648 1658
2755 2798 1405 2704 2738 3008 2586 2873 1232 2597 2516 1537 1128 2849
1399 1131 1569 2381 1084 2304 4530 2297 2279 2303 2669 4225 2513 2725
1955 2527 4443 2478 1919 1813 2533 828 2015 3078 4495 2673 2316 2647
3402 3494 2156 3236 2612 2323 2409 2354 1285 2616 1427 1516 2456 2844
1495 2594 2604 1268 2198 3038 2927]
sqft_lot15 [5650 7639 8062 ... 5731 1509 2007]
sales_year [2014 2015]
sales_month [10 12 2 5 6 1 4 3 7 8 11 9]

```

2.1.3 4c) Drop features [5 pts]

Let's drop features that are unnecessary. `id` is not a meaningful feature. `date` string has been coded to `sales_month` and `sales_year`, so we can remove `date`. `zipcode` can be also removed as it's hard to include in a linear regressio model and the location info is included in the `lat` and `long`. Drop the features `id`, `date`, and `zipcode` and replace the `df2`.

```

[330]: # drop unnecessary features, replace df2
       # your code here
       df2.drop(['id', 'date', 'zipcode'],axis=1, inplace=True)

```

```

[331]: # tests that you droppd the features id, date, and zipcode from df2

```

2.2 5. More inspection; Correlation and pair plot [5 pts and Peer Review]

2.2.1 5a) Get correlation matrix on the data frame. [5 pts]

Which feature may be the best predictor of price based on the correlation? Answer as a string value (e.g. `best_guess_predictor = 'price'` or `best_guess_predictor = 'yr_built'`)

```
[332]: # your code here
print(df2.corr())

# uncomment and update best_guess_predictor with a string value
best_guess_predictor = 'sqft_living'
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	\
price	1.000000	0.308350	0.525138	0.702035	0.089661	0.256794	
bedrooms	0.308350	1.000000	0.515884	0.576671	0.031703	0.175429	
bathrooms	0.525138	0.515884	1.000000	0.754665	0.087740	0.500653	
sqft_living	0.702035	0.576671	0.754665	1.000000	0.172826	0.353949	
sqft_lot	0.089661	0.031703	0.087740	0.172826	1.000000	-0.005201	
floors	0.256794	0.175429	0.500653	0.353949	-0.005201	1.000000	
waterfront	0.266369	-0.006582	0.063744	0.103818	0.021604	0.023698	
view	0.397293	0.079532	0.187737	0.284611	0.074710	0.029444	
condition	0.036362	0.028472	-0.124982	-0.058753	-0.008958	-0.263768	
grade	0.667434	0.356967	0.664983	0.762704	0.113621	0.458183	
sqft_above	0.605567	0.477600	0.685342	0.876597	0.183512	0.523885	
sqft_basement	0.323816	0.303093	0.283770	0.435043	0.015286	-0.245705	
yr_built	0.054012	0.154178	0.506019	0.318049	0.053080	0.489319	
yr_renovated	0.126434	0.018841	0.050739	0.055363	0.007644	0.006338	
lat	0.307003	-0.008931	0.024573	0.052529	-0.085683	0.049614	
long	0.021626	0.129473	0.223042	0.240223	0.229521	0.125419	
sqft_living15	0.585379	0.391638	0.568634	0.756420	0.144608	0.279885	
sqft_lot15	0.082447	0.029244	0.087175	0.183286	0.718557	-0.011269	
sales_year	0.003576	-0.009838	-0.026596	-0.029038	0.005468	-0.022315	
sales_month	-0.010081	-0.001533	0.007392	0.011810	-0.002369	0.014005	

	waterfront	view	condition	grade	sqft_above	\
price	0.266369	0.397293	0.036362	0.667434	0.605567	
bedrooms	-0.006582	0.079532	0.028472	0.356967	0.477600	
bathrooms	0.063744	0.187737	-0.124982	0.664983	0.685342	
sqft_living	0.103818	0.284611	-0.058753	0.762704	0.876597	
sqft_lot	0.021604	0.074710	-0.008958	0.113621	0.183512	
floors	0.023698	0.029444	-0.263768	0.458183	0.523885	
waterfront	1.000000	0.401857	0.016653	0.082775	0.072075	
view	0.401857	1.000000	0.045990	0.251321	0.167649	
condition	0.016653	0.045990	1.000000	-0.144674	-0.158214	
grade	0.082775	0.251321	-0.144674	1.000000	0.755923	
sqft_above	0.072075	0.167649	-0.158214	0.755923	1.000000	
sqft_basement	0.080588	0.276947	0.174105	0.168392	-0.051943	
yr_built	-0.026161	-0.053440	-0.361417	0.446963	0.423898	
yr_renovated	0.092885	0.103917	-0.060618	0.014414	0.023285	
lat	-0.014274	0.006157	-0.014941	0.114084	-0.000816	
long	-0.041910	-0.078400	-0.106500	0.198372	0.343803	
sqft_living15	0.086463	0.280439	-0.092824	0.713202	0.731870	
sqft_lot15	0.030703	0.072575	-0.003406	0.119248	0.194050	

sales_year	-0.004165	0.001364	-0.045589	-0.030387	-0.023823
sales_month	0.008132	-0.005638	0.021978	0.008376	0.009872

	sqft_basement	yr_built	yr_renovated	lat	long \
price	0.323816	0.054012	0.126434	0.307003	0.021626
bedrooms	0.303093	0.154178	0.018841	-0.008931	0.129473
bathrooms	0.283770	0.506019	0.050739	0.024573	0.223042
sqft_living	0.435043	0.318049	0.055363	0.052529	0.240223
sqft_lot	0.015286	0.053080	0.007644	-0.085683	0.229521
floors	-0.245705	0.489319	0.006338	0.049614	0.125419
waterfront	0.080588	-0.026161	0.092885	-0.014274	-0.041910
view	0.276947	-0.053440	0.103917	0.006157	-0.078400
condition	0.174105	-0.361417	-0.060618	-0.014941	-0.106500
grade	0.168392	0.446963	0.014414	0.114084	0.198372
sqft_above	-0.051943	0.423898	0.023285	-0.000816	0.343803
sqft_basement	1.000000	-0.133124	0.071323	0.110538	-0.144765
yr_built	-0.133124	1.000000	-0.224874	-0.148122	0.409356
yr_renovated	0.071323	-0.224874	1.000000	0.029398	-0.068372
lat	0.110538	-0.148122	0.029398	1.000000	-0.135512
long	-0.144765	0.409356	-0.068372	-0.135512	1.000000
sqft_living15	0.200355	0.326229	-0.002673	0.048858	0.334605
sqft_lot15	0.017276	0.070958	0.007854	-0.086419	0.254451
sales_year	-0.015687	0.003507	-0.023707	-0.029212	0.000270
sales_month	0.006035	-0.006226	0.012827	0.014961	-0.008134

	sqft_living15	sqft_lot15	sales_year	sales_month
price	0.585379	0.082447	0.003576	-0.010081
bedrooms	0.391638	0.029244	-0.009838	-0.001533
bathrooms	0.568634	0.087175	-0.026596	0.007392
sqft_living	0.756420	0.183286	-0.029038	0.011810
sqft_lot	0.144608	0.718557	0.005468	-0.002369
floors	0.279885	-0.011269	-0.022315	0.014005
waterfront	0.086463	0.030703	-0.004165	0.008132
view	0.280439	0.072575	0.001364	-0.005638
condition	-0.092824	-0.003406	-0.045589	0.021978
grade	0.713202	0.119248	-0.030387	0.008376
sqft_above	0.731870	0.194050	-0.023823	0.009872
sqft_basement	0.200355	0.017276	-0.015687	0.006035
yr_built	0.326229	0.070958	0.003507	-0.006226
yr_renovated	-0.002673	0.007854	-0.023707	0.012827
lat	0.048858	-0.086419	-0.029212	0.014961
long	0.334605	0.254451	0.000270	-0.008134
sqft_living15	1.000000	0.183192	-0.021734	0.002449
sqft_lot15	0.183192	1.000000	-0.000085	0.003546
sales_year	-0.021734	-0.000085	1.000000	-0.782389
sales_month	0.002449	0.003546	-0.782389	1.000000

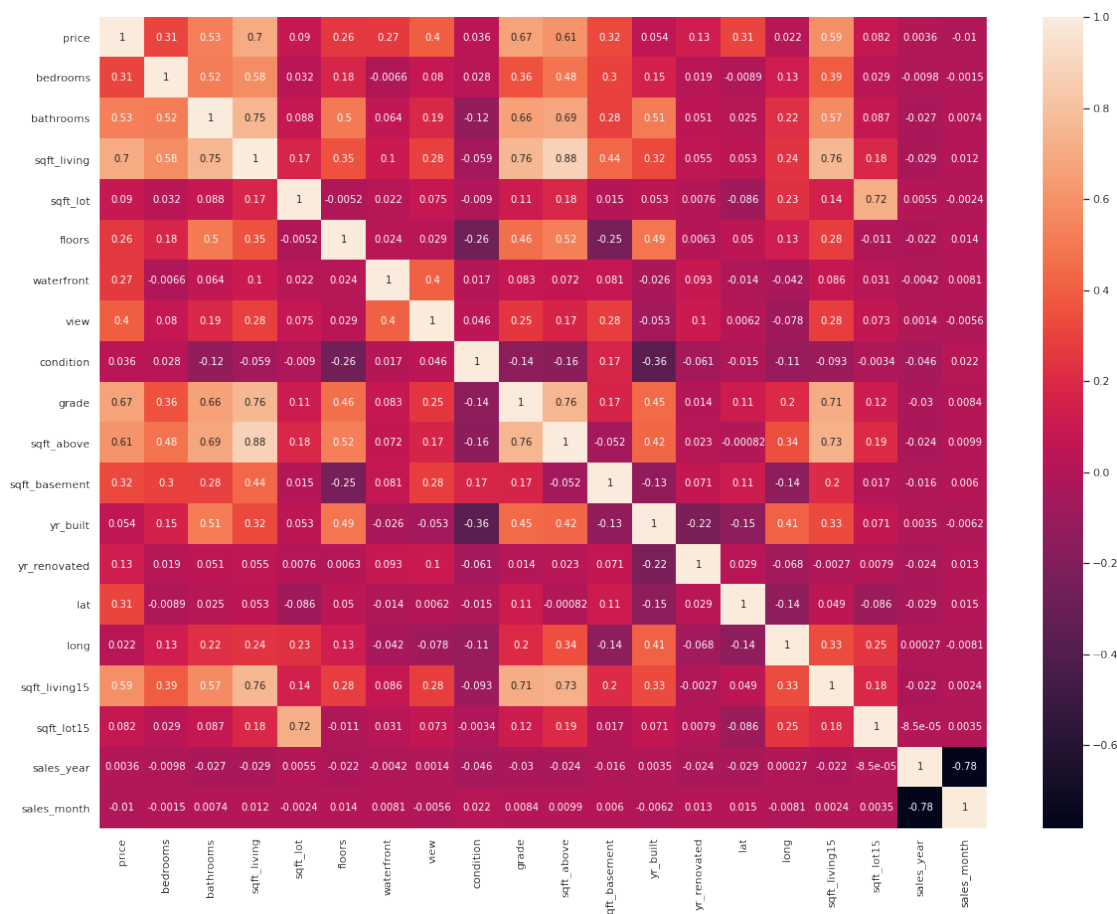
```
[333]: # tests the solution for best_guess_predictor
```

2.2.2 5b) Display the correlation matrix as heat map [Peer Review]

`seaborn.heatmap()` can visualize a matrix as a heatmap. Visualize the correlation matrix using `seaborn.heatmap()`. Play with color map, text font size, decimals, text orientation etc. If you find how to make a pretty visualization, please share in the discussion board. You will upload your correlation matrix in the Peer Review assignment for the week. **Note:** your code for this section may cause the Validate button to time out. If you want to run the Validate button prior to submitting, you could comment out the code in this section after completing the Peer Review.

```
[334]: # practice visualizing correlation matrix using a heatmap
# your code here
plt.subplots(figsize=(20,15))
sns.heatmap(df2.corr(), annot=True)
```

```
[334]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1f1bf4bad0>
```



2.2.3 5c) Pair plot [Peer Review]

Pair plot is a fast way to inspect relationships between features. Use seaborn's `.pairplot()` function to draw a pairplot if the first 10 columns (including price) and inspect their relationships. Set the diagonal elements to be KDE plot. You will upload your pair plot in this week's Peer Review assignment. **Note:** your code for this section may cause the Validate button to time out. If you want to run the Validate button prior to submitting, you could comment out the code in this section after completing the Peer Review.

```
[335]: # practice inspecting relationships between features using a pair plot.  
# your code here  
sns.pairplot(df2.iloc[:, :10], diag_kind='kde')
```

[335]: <seaborn.axisgrid.PairGrid at 0x7f1f1bf529d0>



2.3 6. Simple linear regression [Peer Review]

2.3.1 6a) Data preparation [Peer Review]

We will split the data to train and test datasets such that the test dataset is 20% of original data. Use `sklearn.model_selection.train_test_split` function to split the data frame to `X_train` and `X_test`. `X_train` is 80% of observation randomly chosen. `X_test` is the rest 20%. Both `X_train` and `X_test` are `pd.DataFrame` object and include 'price' in the table. Note that the `train_test_split` can handle data frame as well as array.

```
[336]: # your code here
from sklearn.model_selection import train_test_split as tst
X_train, X_test = tst(df2, test_size=0.20, random_state=15413)

# use sklearn.model_selection.train_test_split to split the data frame
# X_train is 80% of the observations; X_test is 20% of the observations
# print length of X_train and X_test
print(len(X_train))
print(len(X_test))
```

17290
4323

```
[337]: # instructor testing cell
# your code here
```

2.3.2 6b) Train a simple linear regression model [Peer Review]

Use the `best_guess_predictor` as a single predictor and build a simple linear regression model using `statsmodels.formula.api.ols` function (https://www.statsmodels.org/dev/example_formulas.html) Print out the result summary. Train on the `X_train` portion. What is the adjusted R-squared value?

```
[338]: # use best_guess_predictor as a single predictor
# build a simple linear regression model, train on the X_train portion
# your code here

model = smf.ols(formula='price ~ sqft_living', data=X_train)
res = model.fit()
print(res.summary())

adj_R2 = 0.493 #update this value according to the result
```

OLS Regression Results

=====

```

Dep. Variable:          price    R-squared:                0.493
Model:                  OLS      Adj. R-squared:           0.493
Method:                 Least Squares    F-statistic:             1.684e+04
Date:                   Sat, 16 Apr 2022    Prob (F-statistic):       0.00
Time:                   01:33:25    Log-Likelihood:          -2.4028e+05
No. Observations:      17290    AIC:                     4.806e+05
Df Residuals:          17288    BIC:                     4.806e+05
Df Model:               1
Covariance Type:        nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept    -4.468e+04    4923.238     -9.075     0.000    -5.43e+04    -3.5e+04
sqft_living    281.1457        2.167    129.764     0.000     276.899     285.392
=====
Omnibus:            11584.466    Durbin-Watson:           1.997
Prob(Omnibus):      0.000    Jarque-Bera (JB):        380472.155
Skew:               2.760    Prob(JB):                 0.00
Kurtosis:           25.308    Cond. No.                 5.60e+03
=====

```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.6e+03. This might indicate that there are strong multicollinearity or other numerical problems.

2.3.3 6c) Best predictor [Peer Review]

In question 5a, we picked a best guess predictor for price based on the correlation matrix. Now we will consider whether the best_guess_predictor that we used is still the best. Print out a list ranking all of the predictors. Then print out a list of the top three predictors in order.

Hint: Linear regression uses adjusted R squared as fit performance. In this week's Peer Review, answer the following questions: What were your top three predictors? How did you order your list of predictors to select those as the top ones? Is your top predictor for this section the same as the best guess predictor you selected in question 5a?

```

[339]: # your code here
      ### Ranking og each predictor and their adjusted R^2

      # sqft_living - 0.493
      # grade - 0.450
      # sqft_above - 0.371
      # sqft_living15 - 0.346
      # bathrooms - 0.281
      # view - 0.16
      # sqft_basement - 0.102

```



```
# bedrooms - 0.098
# lat - 0.094
# floors - 0.066
# waterfront - 0.066
# yr_renovated - 0.016
# sqft_lot - 0.008
# sqft_lot15 - 0.008
# yr_built - 0.003
# condition - 0.001
# long - 0
# sales_year - 0
# sales_month - 0

# uncomment and update top_three
top_three = ['sqft_living', 'grade', 'sqft_above']
```

```
[340]: # instructor testing cell
# your code here
```