

```
In [ ]: from google.colab import files
upload = files.upload()
```

Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving mushrooms.csv to mushrooms.csv

```
In [ ]: import pandas as pd
import numpy as np
import seaborn as sns
```

```
In [ ]: mushroom = pd.read_csv('mushrooms.csv')
mushroom.head()
```

```
Out [ ]:
```

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	stalk-shape	stalk-root	stalk-surface-above-ring
0	p	x	s	n	t	p	f	c	n	k	e	e	s
1	e	x	s	y	t	a	f	c	b	k	e	c	s
2	e	b	s	w	t	l	f	c	b	n	e	c	s
3	p	x	y	w	t	p	f	c	n	n	e	e	s
4	e	x	s	g	f	n	f	w	b	k	t	e	s

```
In [ ]: target = mushroom['class']
def fix(x):
    if x=='p':
        return 1
    else:
        return 0
target = target.apply(fix)

features = mushroom.iloc[:, 1:]
features = pd.get_dummies(features, prefix = features.columns, drop_first=True)
print('The number of entries is:', len(features))
print('The proportion of poisonous entries is:', sum(target)/len(target))
```

The number of entries is: 8124

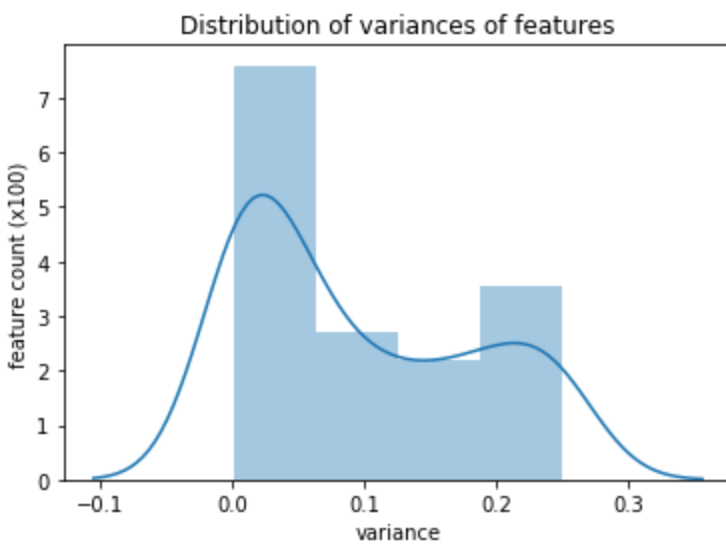
The proportion of poisonous entries is: 0.48202855736090594

```
In [ ]: # variance thesholding
col_sum = features.apply(np.var, axis=0)
col_sort = col_sum.sort_values()
print('The number of expanded features is:', len(col_sum))
print('The number of features with variance less than 0.05:', len(col_sort[col_sort<0.1])
      # number of columns that define less than 3% of the data

# plot of variances, cut off at 0.1
col_sort
g = sns.distplot(col_sort)
g.set_title('Distribution of variances of features')
g.set_ylabel('feature count (x100)')
g.set_xlabel('variance');
```

The number of expanded features is: 95

The number of features with variance less than 0.05: 56



```
In [ ]: columns = col_sort[col_sort>=0.1]
print('The final number of features to be used is:', len(columns))
features_final = features[columns.index]
```

The final number of features to be used is: 39

```
In [ ]: # run to silence unnecessary warnings
def warn(*args, **kwargs):
    pass
import warnings
warnings.warn = warn
```

```
In [ ]: # extensive gridsearch
def perform_gridsearch(features, target):
    '''Performs gridsearch on Logistic Regression'''
    import numpy as np
    from sklearn.linear_model import LogisticRegression
    from sklearn.model_selection import GridSearchCV

    # # # Create model (can change to a different model/parameters)
    model = LogisticRegression()
    # Create range of candidate penalty hyperparameter values
    penalty = ['l1', 'l2']
    # Create range of candidate regularization hyperparameter values C
    C = np.logspace(0, 4, 10)
    # Create dictionary hyperparameter candidates
    hyperparameters = dict(C=C, penalty=penalty)

    # Create grid search, and pass in all defined values
    gridsearch = GridSearchCV(model, hyperparameters, cv=5, verbose=1) # Fit grid search
    best_model = gridsearch.fit(features, target)
    return best_model
    print('Best Penalty:', best_model.best_estimator_.get_params()['penalty'])
    print('Best C:', best_model.best_estimator_.get_params()['C'])
    print("The mean accuracy of the model is:", best_model.score(features, target))
```

```
In [ ]: grid = perform_gridsearch(features_final, target)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 13.6s finished

```
In [ ]: print(grid.best_estimator_.get_params())
print(grid.score(features_final, target))
```

{'C': 1.0, 'class_weight': None, 'dual': False, 'fit_intercept': True, 'intercept_scaling': 1, 'l1_penalty': 0.0, 'l2_penalty': 0.0, 'max_iter': 1000, 'multi_class': 'ovr', 'n_jobs': 1, 'penalty': 'l2', 'random_state': None, 'solver': 'lbfgs', 'tol': 0.0001, 'verbose': 0, 'warm_start': False}

```
g': 1, 'l1_ratio': None, 'max_iter': 100, 'multi_class': 'warn', 'n_jobs': None, 'penalty': 'l1', 'random_state': None, 'solver': 'warn', 'tol': 0.0001, 'verbose': 0, 'warm_start': False}
0.9988921713441654
```

```
In [ ]: # randomized search
def perform_randomized_search(features, target):
    '''Performs a randomized search on Logistic regression'''
    import numpy as np
    from sklearn.linear_model import LogisticRegression
    from sklearn.model_selection import RandomizedSearchCV
    from scipy.stats import uniform

    # Create logistic regression
    model = LogisticRegression()
    # Create range of candidate regularization penalty hyperparameter values
    penalty = ['l1', 'l2']
    # Create distribution of candidate regularization hyperparameter values
    C = uniform(loc=0, scale=4)
    # Create hyperparameter options
    hyperparameters = dict(C=C, penalty=penalty)

    # Create randomized search
    randomizedsearch = RandomizedSearchCV(
        model, hyperparameters, random_state=1, n_iter=100, cv=5, verbose=1)
    # Fit randomized search
    best_model = randomizedsearch.fit(features, target)
    return best_model
    print(best_model.best_estimator_)
    print("The mean accuracy of the model is:",best_model.score(features, target))
```

```
In [ ]: random = perform_randomized_search(features_final, target)

Fitting 5 folds for each of 100 candidates, totalling 500 fits
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 500 out of 500 | elapsed: 54.5s finished
```

```
In [ ]: print(random.best_estimator_.get_params())
print(random.score(features_final, target))

{'C': 0.9183088549193021, 'class_weight': None, 'dual': False, 'fit_intercept': True, 'intercept_scaling': 1, 'l1_ratio': None, 'max_iter': 100, 'multi_class': 'warn', 'n_jobs': None, 'penalty': 'l1', 'random_state': None, 'solver': 'warn', 'tol': 0.0001, 'verbose': 0, 'warm_start': False}
0.9988921713441654
```

```
In [ ]: # pipelines with grid search
def execute_pipeline(features,target):
    '''
    Takes in features(X) and target(y) and performs a grid search classification
    by Logisitic Regression, SVM, and Random Forest. Prints the best model.
    '''
    import numpy as np
    from scipy.stats import uniform
    from sklearn.linear_model import LogisticRegression
    from sklearn.model_selection import GridSearchCV
    from sklearn.pipeline import Pipeline, FeatureUnion
    from sklearn.decomposition import PCA
    from sklearn.preprocessing import StandardScaler
    from sklearn.svm import SVC
    from sklearn.ensemble import RandomForestClassifier

    # C = uniform(loc=0, scale=4) # optional space of candidate values
```

```

# preprocessing and creating the pipe
pca_components = PCA() ## if n_components not specified, keeps all components
std_scaler = StandardScaler()
preprocess = FeatureUnion([("std",std_scaler), ("pca", pca_components)])
model = LogisticRegression()
pipe = Pipeline([("preprocess", preprocess),
                  ("classifier", model)])

# can add in other models here, just add to search_space below
log_reg = {"classifier": [LogisticRegression()],
           "classifier__penalty": ['l2','l1'],
           "classifier__C": np.logspace(0, 4, 10)}
svm = {"classifier": [SVC()],
       "classifier__C": [0.01,0.1,1,10,100],
       "classifier__kernel":['linear','poly','rbf','sigmoid']}
log_reg_solver = {"classifier": [LogisticRegression()],
                  "classifier__penalty": ['l2'],
                  "classifier__C": np.logspace(0, 4, 10),
                  "classifier__solver":['newton-cg','saga','sag','liblinear']} ##These so
rand_forest = {"classifier": [RandomForestClassifier()],
               "classifier__n_estimators": [10, 100, 1000],
               "classifier__max_depth": [5,8,15,25,30, None],
               "classifier__min_samples_leaf": [1,2,5,10,15,100],
               "classifier__max_leaf_nodes": [2, 5,10]}
search_space = [svm,log_reg, log_reg_solver, rand_forest]

# executing pipeline
clf = GridSearchCV(pipe, search_space, cv=5, verbose=1, n_jobs=-1) # Fit grid search
best_model = clf.fit(features, target)
return best_model
print(best_model.best_estimator_)
print("The mean accuracy of the model is:",best_model.score(features, target))

```

```
In [ ]: big = execute_pipeline(features_final, target)
```

Fitting 5 folds for each of 404 candidates, totalling 2020 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 46 tasks      | elapsed: 1.2min
[Parallel(n_jobs=-1)]: Done 196 tasks    | elapsed: 2.2min
[Parallel(n_jobs=-1)]: Done 446 tasks    | elapsed: 4.9min
[Parallel(n_jobs=-1)]: Done 796 tasks    | elapsed: 16.8min
[Parallel(n_jobs=-1)]: Done 1246 tasks   | elapsed: 33.5min
[Parallel(n_jobs=-1)]: Done 1796 tasks   | elapsed: 53.3min
[Parallel(n_jobs=-1)]: Done 2020 out of 2020 | elapsed: 62.0min finished

```

```
In [ ]: print(big.best_estimator_[ 'classifier' ])
print(big.score(features_final, target))

SVC(C=10, cache_size=200, class_weight=None, coef0=0.0,
     decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
     kernel='linear', max_iter=-1, probability=False, random_state=None,
     shrinking=True, tol=0.001, verbose=False)
1.0
```

```
In [ ]: # visualizing effects of number of trees on random forest
import matplotlib.pyplot as plt
import numpy as np
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import validation_curve

param_range = np.arange(1, 250, 2)
# Calculate accuracy on training and test set using range of parameter values

train_scores, test_scores = validation_curve(
    # Classifier

```

```

RandomForestClassifier(),
# Feature matrix features,
X=features_final,
# Target vector
y=target,
# Hyperparameter to examine
param_name="n_estimators",
# Range of hyperparameter's values
param_range=param_range,
# Number of folds
cv=3,
# Performance metric
scoring="accuracy",
# Use all computer cores
n_jobs=-1)

# Calculate mean and standard deviation for training set scores
train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)
# Calculate mean and standard deviation for test set scores
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)

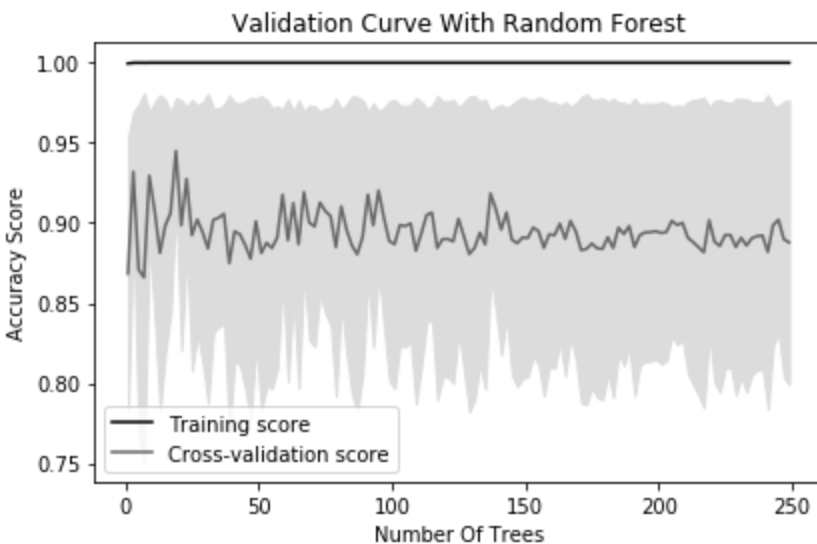
```

```

In [ ]: # Plot mean accuracy scores for training and test sets
# # # #
plt.plot(param_range, train_mean, label="Training score", color="black")
plt.plot(param_range, test_mean, label="Cross-validation score", color="dimgrey")
# # # #
# Plot accuracy bands for training and test sets
plt.fill_between(param_range, test_mean - test_std,
                 test_mean + test_std, color="gainsboro")

# Create plot
plt.title("Validation Curve With Random Forest")
plt.xlabel("Number Of Trees")
plt.ylabel("Accuracy Score")
plt.tight_layout()
plt.legend(loc="best")
#plt.savefig('RF.png')

```



```

In [ ]: # Load libraries
import matplotlib.pyplot as plt
import numpy as np
from sklearn.svm import SVC
from sklearn.model_selection import validation_curve

param_range = np.arange(0.1, 10, 0.1)

```

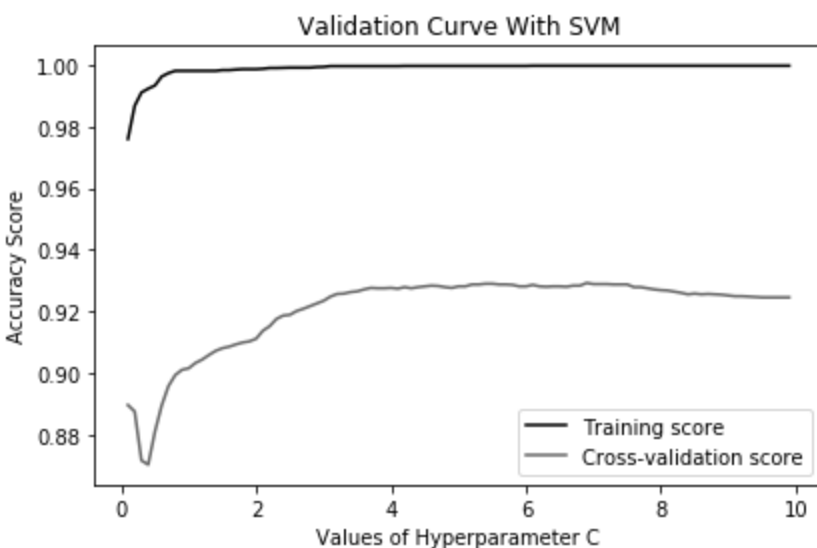
```
# Calculate accuracy on training and test set using range of parameter values
```

```
train_scores, test_scores = validation_curve( # Classifier
                                             SVC(),
                                             # Feature matrix features,
                                             X=features_final,
                                             # Target vector
                                             y=target,
                                             # Hyperparameter to examine
                                             param_name="C",
                                             # Range of hyperparameter's values
                                             param_range=param_range,
                                             # Number of folds
                                             cv=3,
                                             # Performance metric
                                             scoring="accuracy",
                                             n_jobs=-1)

# Calculate mean and standard deviation for training set scores
train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)
# Calculate mean and standard deviation for test set scores
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)
```

```
In [ ]: # Plot mean accuracy scores for training and test sets
        #####
        plt.plot(param_range, train_mean, label="Training score", color="black")
        plt.plot(param_range, test_mean, label="Cross-validation score", color="dimgrey")
        #####

        # Create plot
        plt.title("Validation Curve With SVM")
        plt.xlabel("Values of Hyperparameter C")
        plt.ylabel("Accuracy Score")
        plt.tight_layout()
        plt.legend(loc="best")
        #plt.savefig('SVM.png')
```



```
In [ ]: # Load libraries
import matplotlib.pyplot as plt
import numpy as np

from sklearn.model_selection import validation_curve
```

```

param_range = np.arange(0.05, 2, 0.05)
# Calculate accuracy on training and test set using range of parameter values

train_scores, test_scores = validation_curve( # Classifier
                                             LogisticRegression(),
                                             # Feature matrix features,
                                             X=features_final,
                                             # Target vector
                                             y=target,
                                             # Hyperparameter to examine
                                             param_name="C",
                                             # Range of hyperparameter's values
                                             param_range=param_range,
                                             # Number of folds
                                             cv=3,
                                             # Performance metric
                                             scoring="accuracy")

# Calculate mean and standard deviation for training set scores
train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)
# Calculate mean and standard deviation for test set scores
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)

```

```

In [ ]: # Plot mean accuracy scores for training and test sets
# # # #
plt.plot(param_range, train_mean, label="Training score", color="black")
plt.plot(param_range, test_mean, label="Cross-validation score", color="dimgrey")
# # # #

# Create plot
plt.title("Validation Curve With Logistic Regression")
plt.xlabel("Values of Hyperparameter C")
plt.ylabel("Accuracy Score")
plt.tight_layout()
plt.legend(loc="best")
#plt.savefig('LR.png')

```

