

# Chapter 03 - Classification - 2

March 27, 2022

## 1 Classification (02)

In this lab we would be going through: - Naive Bayes - K-Nearest Neighbours - Poisson Regression

For this lab, we would examine the **Smarket** data set that contains a number of numeric variables plus a variable called **Direction** which has the two labels **Up** and **Down**

Our goal is to predict **Direction** using the other features

```
[1]: library(e1071)
      library(ISLR2)
      attach(Smarket)

[2]: train <- (Year < 2005)

      # Test data
      Smarket.test <- Smarket[!train, ]
      dim(Smarket.test)

      #Train data
      Smarket.train = Smarket[train, ]
      dim(Smarket.train)

      Direction.2005 = Direction[!train]
```

1. 252 2. 9

1. 998 2. 9

### 1.1 Naive Bayes

We are using the **naiveBayes()** function, which is part of the **e1071** **naiveBayes()** library.

By default, this implementation of the naive Bayes classifier models each quantitative feature using a Gaussian distribution. However, a kernel density method can also be used to estimate the distributions.

```
[3]: nb.fit <- naiveBayes(Direction ~ Lag1 + Lag2, data = Smarket,
      subset = train)
```

```
nb.fit
```

Naive Bayes Classifier for Discrete Predictors

Call:

```
naiveBayes.default(x = X, y = Y, laplace = laplace)
```

A-priori probabilities:

```
Y
  Down      Up
0.491984 0.508016
```

Conditional probabilities:

```
      Lag1
Y      [,1] [,2]
Down 0.04279022 1.227446
Up   -0.03954635 1.231668
```

```
      Lag2
Y      [,1] [,2]
Down 0.03389409 1.239191
Up   -0.03132544 1.220765
```

The output contains the estimated mean and standard deviation for each variable in each class.

```
[4]: mean(Lag1[train][Direction[train] == 'Down'])
      sd(Lag1[train][Direction[train] == 'Down'])
```

```
0.0427902240325866
```

```
1.22744562820108
```

```
[5]: nb.class = predict(nb.fit, Smarket.test)
      table(nb.class, Direction.2005)

      mean(nb.class == Direction.2005)
```

```
      Direction.2005
nb.class Down  Up
Down    28  20
Up     83 121
```

```
0.591269841269841
```

Naive Bayes performs very well on this data, with accurate predictions over 59% of the time. This is slightly worse than QDA, but much better than LDA.

The `predict()` function can also generate estimates of the probability that each observation belongs to a particular class

```
[6]: nb.preds = predict(nb.fit, Smarket.test, type = "raw")
      nb.preds[1:5, ]
```

A matrix: 5 × 2 of type dbl

	Down	Up
	0.4873164	0.5126836
	0.4762492	0.5237508
	0.4653377	0.5346623
	0.4748652	0.5251348
	0.4901890	0.5098110

## 1.2 K - Nearest Neighbors

We would be using the `knn()` function which is a part of the `class` library. Rather than a two-step approach in which we first fit the model and then we use the model to make predictions, `knn()` forms predictions using a single command.

The function requires four inputs: 1. A matrix containing the predictors associated with the training data, labeled `train.X` below. 2. A matrix containing the predictors associated with the data for which we wish to make predictions, labeled `test.X` below. 3. A vector containing the class labels for the training observations, labeled `train.Direction` below. 4. A value for K, the number of nearest neighbors to be used by the classifier.

```
[7]: library(class)
```

```
[8]: train.X = cbind(Lag1, Lag2)[train, ] #cbind() is short for column bind, binds
      ↪ variables together
      test.X = cbind(Lag1, Lag2)[!train, ]
      train.Direction = Direction[train]
```

We set a random seed before we apply `knn()` because if several observations are tied as nearest neighbors, then R will randomly break the tie.

```
[9]: set.seed(1)
      knn.pred = knn(train.X, test.X, train.Direction, k=1)

      table(knn.pred, Direction.2005)
      mean(knn.pred==Direction.2005) #performance
```

	Direction.2005	
knn.pred	Down	Up
Down	43	58
Up	68	83

0.5

The results using  $K = 1$  are not very good, since only 50 % of the observations are correctly predicted. Of course, it may be that  $K = 1$  results in an overly flexible fit to the data.

```
[12]: #return a k-nn model with three neighbors
knn.pred = function(){
  # your code here
  return(knn(train.X, test.X, train.Direction, k=3))
}
knn.pred = knn.pred()
```

```
[13]: table(knn.pred, Direction.2005)

#Test the performance of new model
stopifnot(round(mean(knn.pred == Direction.2005),2) == 0.54)
```

	Direction.2005	
knn.pred	Down	Up
Down	48	54
Up	63	87

```
[14]: knn.pred = knn(train.X, test.X, train.Direction, k=4)
mean(knn.pred == Direction.2005)
```

0.496031746031746

We can see that the results have improved slightly when we increase the value of  $K$  from 1 to 3. But increasing  $K$  further turns out to provide no further improvements.

It appears that for this data, QDA provides the best results of the methods that we have examined so far.

### 1.3 Poisson Regression

We would be using the `glm()` function with the argument `family = poisson` to define a poisson regression model.

We are gonna fit a Poisson regression model to the **Bikeshare** data set found in **ISLR2** library, which measures the number of bike rentals(**bikers**) per hour in Washington DC.

```
[15]: attach(Bikeshare) #attaching the data set to R's context
```

```
[16]: dim(Bikeshare)
names(Bikeshare)
```

1. 8645 2. 15

1. 'season' 2. 'mnth' 3. 'day' 4. 'hr' 5. 'holiday' 6. 'weekday' 7. 'workingday' 8. 'weathersit' 9. 'temp' 10. 'atemp' 11. 'hum' 12. 'windspeed' 13. 'casual' 14. 'registered' 15. 'bikers'

```
[17]: mod.pois = glm(bikers ~ mnth + hr + workingday + temp + weathersit,
                  data = Bikeshare, family = poisson)
summary(mod.pois)
```

Call:

```
glm(formula = bikers ~ mnth + hr + workingday + temp + weathersit,
    family = poisson, data = Bikeshare)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-20.7574	-3.3441	-0.6549	2.6999	21.9628

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	2.693688	0.009720	277.124	< 2e-16 ***
mnthFeb	0.226046	0.006951	32.521	< 2e-16 ***
mnthMarch	0.376437	0.006691	56.263	< 2e-16 ***
mnthApril	0.691693	0.006987	98.996	< 2e-16 ***
mnthMay	0.910641	0.007436	122.469	< 2e-16 ***
mnthJune	0.893405	0.008242	108.402	< 2e-16 ***
mnthJuly	0.773787	0.008806	87.874	< 2e-16 ***
mnthAug	0.821341	0.008332	98.573	< 2e-16 ***
mnthSept	0.903663	0.007621	118.578	< 2e-16 ***
mnthOct	0.937743	0.006744	139.054	< 2e-16 ***
mnthNov	0.820433	0.006494	126.334	< 2e-16 ***
mnthDec	0.686850	0.006317	108.724	< 2e-16 ***
hr1	-0.471593	0.012999	-36.278	< 2e-16 ***
hr2	-0.808761	0.014646	-55.220	< 2e-16 ***
hr3	-1.443918	0.018843	-76.631	< 2e-16 ***
hr4	-2.076098	0.024796	-83.728	< 2e-16 ***
hr5	-1.060271	0.016075	-65.957	< 2e-16 ***
hr6	0.324498	0.010610	30.585	< 2e-16 ***
hr7	1.329567	0.009056	146.822	< 2e-16 ***
hr8	1.831313	0.008653	211.630	< 2e-16 ***
hr9	1.336155	0.009016	148.191	< 2e-16 ***
hr10	1.091238	0.009261	117.831	< 2e-16 ***
hr11	1.248507	0.009093	137.304	< 2e-16 ***
hr12	1.434028	0.008936	160.486	< 2e-16 ***
hr13	1.427951	0.008951	159.529	< 2e-16 ***
hr14	1.379296	0.008999	153.266	< 2e-16 ***
hr15	1.408149	0.008977	156.862	< 2e-16 ***
hr16	1.628688	0.008805	184.979	< 2e-16 ***
hr17	2.049021	0.008565	239.221	< 2e-16 ***
hr18	1.966668	0.008586	229.065	< 2e-16 ***
hr19	1.668409	0.008743	190.830	< 2e-16 ***
hr20	1.370588	0.008973	152.737	< 2e-16 ***
hr21	1.118568	0.009215	121.383	< 2e-16 ***

```

hr22                0.871879    0.009536    91.429 < 2e-16 ***
hr23                0.481387    0.010207    47.164 < 2e-16 ***
workingday          0.014665    0.001955     7.502 6.27e-14 ***
temp               0.785292    0.011475    68.434 < 2e-16 ***
weathersitcloudy/misty -0.075231  0.002179   -34.528 < 2e-16 ***
weathersitlight rain/snow -0.575800  0.004058  -141.905 < 2e-16 ***
weathersitheavy rain/snow -0.926287  0.166782   -5.554 2.79e-08 ***

```

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 1052921 on 8644 degrees of freedom  
Residual deviance: 228041 on 8605 degrees of freedom  
AIC: 281159

Number of Fisher Scoring iterations: 5

We are gonna plot these coefficients associated with `mnth` and `hr` for better visualization

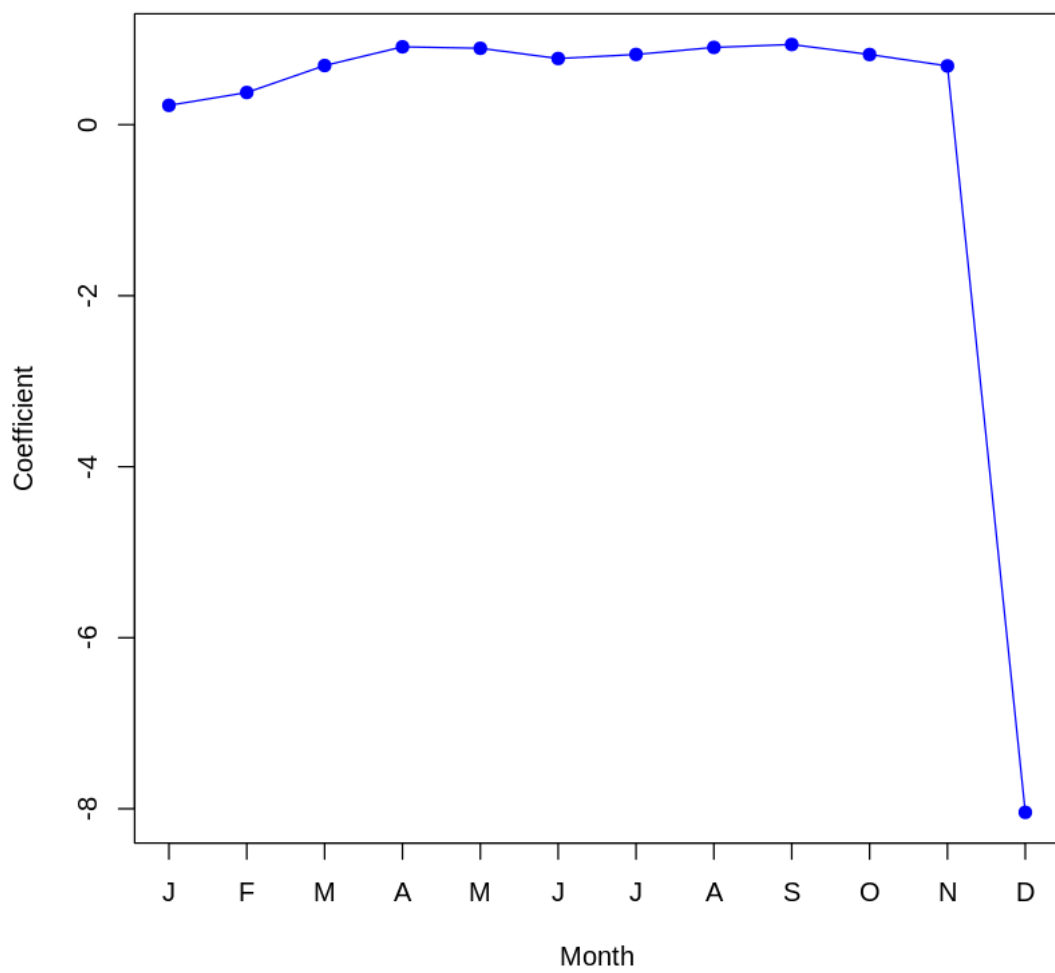
```

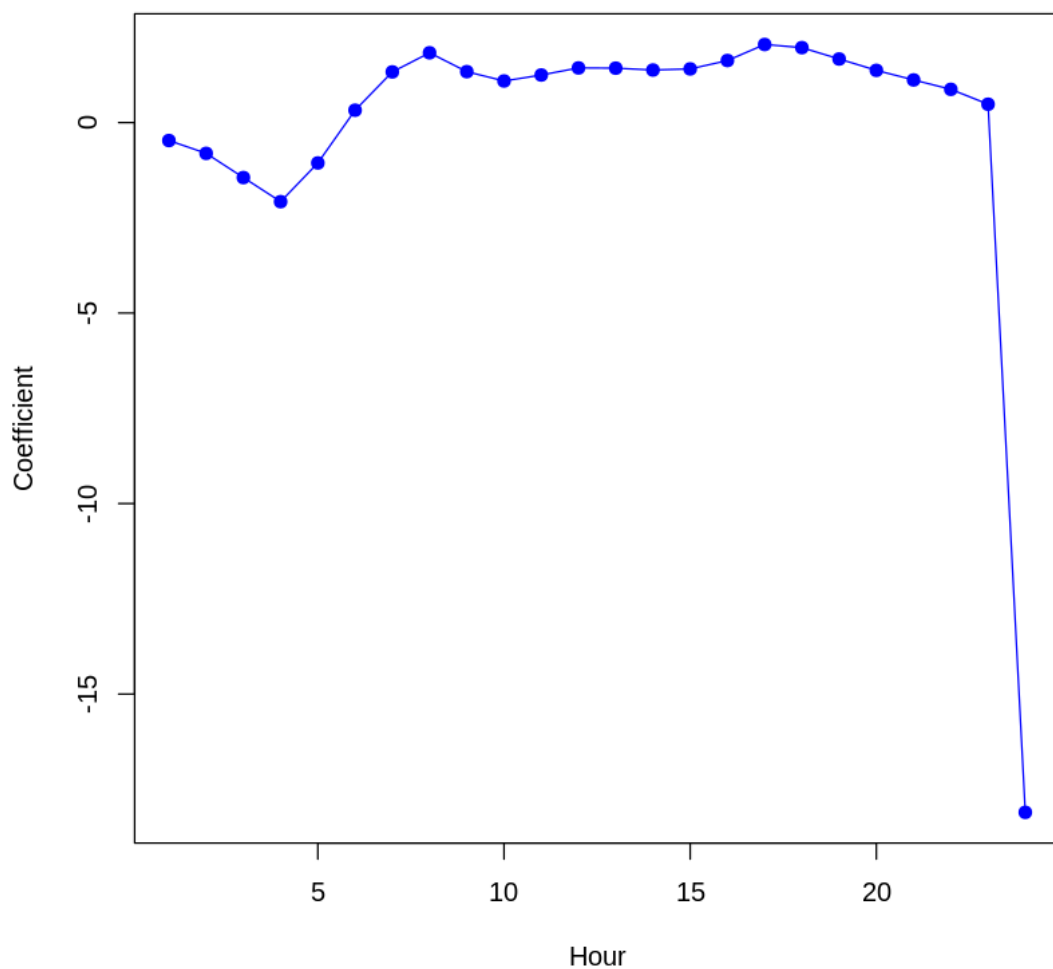
[18]: coef.mnth <- c(coef(mod.pois)[2:12], -sum(coef(mod.pois)[2:12]))

plot(coef.mnth, xlab = "Month", ylab = "Coefficient",
     xaxt = "n", col = "blue", pch = 19, type = "o")
axis(side = 1, at = 1:12,
     labels = c("J", "F", "M", "A", "M", "J", "J", "A", "S", "O", "N", "D"))

coef.hours <- c(coef(mod.pois)[13:35], -sum(coef(mod.pois)[13:35]))
plot(coef.hours, xlab = "Hour", ylab = "Coefficient", col = "blue", pch = 19,
     type = "o")

```





We can once again use the `predict()` function to obtain the fitted values (predictions) from this Poisson regression model.

```
[19]: mod.pred = predict(mod.pois, type = "response")
      summary(mod.pred)
```

```
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
1.201   44.513  124.299  143.794  219.268  585.958
```

```
[ ]:
```