

FinalQuiz

March 28, 2022

1 Final Quiz

```
[1]: library(ISLR2)
      library(MASS)
      library(e1071)
```

Attaching package: ‘MASS’

The following object is masked from ‘package:ISLR2’:

Boston

```
[2]: attach(Auto)
      head(Auto)
```

A data.frame: 6 × 9

	mpg <dbl>	cylinders <int>	displacement <dbl>	horsepower <int>	weight <int>	acceleration <dbl>	year <int>	origin <int>
1	18	8	307	130	3504	12.0	70	1
2	15	8	350	165	3693	11.5	70	1
3	18	8	318	150	3436	11.0	70	1
4	16	8	304	150	3433	12.0	70	1
5	17	8	302	140	3449	10.5	70	1
6	15	8	429	198	4341	10.0	70	1

- (a) Use the `lm()` function to perform a multiple linear regression with `mpg` as the response and all other variables except `name` as the predictors

```
[3]: #lm.fit = ?
      # your code here
      lm.fit = lm(mpg ~ cylinders + displacement + horsepower + weight + acceleration_
      ↪+ year + origin, data=Auto)
      summary(lm.fit)
```

Call:

```
lm(formula = mpg ~ cylinders + displacement + horsepower + weight +
    acceleration + year + origin, data = Auto)
```

Residuals:

Min	1Q	Median	3Q	Max
-9.5903	-2.1565	-0.1169	1.8690	13.0604

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-17.218435	4.644294	-3.707	0.00024	***
cylinders	-0.493376	0.323282	-1.526	0.12780	
displacement	0.019896	0.007515	2.647	0.00844	**
horsepower	-0.016951	0.013787	-1.230	0.21963	
weight	-0.006474	0.000652	-9.929	< 2e-16	***
acceleration	0.080576	0.098845	0.815	0.41548	
year	0.750773	0.050973	14.729	< 2e-16	***
origin	1.426141	0.278136	5.127	4.67e-07	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.328 on 384 degrees of freedom

Multiple R-squared: 0.8215, Adjusted R-squared: 0.8182

F-statistic: 252.4 on 7 and 384 DF, p-value: < 2.2e-16

[4]: *#hidden test cases*

(b) Add an interaction term between `horsepower` and `weight` to the set of existant predictors and define a fit

```
[5]: #lm.fit = ?
# your code here
lm.fit = lm(lm(mpg ~ cylinders + displacement + horsepower + weight +
    ↪acceleration + year + origin + horsepower:weight
    , data=Auto))
summary(lm.fit)
```

Call:

```
lm(formula = lm(mpg ~ cylinders + displacement + horsepower +
    weight + acceleration + year + origin + horsepower:weight,
    data = Auto))
```

Residuals:

Min	1Q	Median	3Q	Max
-8.589	-1.617	-0.184	1.541	12.001

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.876e+00	4.511e+00	0.638	0.524147
cylinders	-2.955e-02	2.881e-01	-0.103	0.918363
displacement	5.950e-03	6.750e-03	0.881	0.378610
horsepower	-2.313e-01	2.363e-02	-9.791	< 2e-16 ***
weight	-1.121e-02	7.285e-04	-15.393	< 2e-16 ***
acceleration	-9.019e-02	8.855e-02	-1.019	0.309081
year	7.695e-01	4.494e-02	17.124	< 2e-16 ***
origin	8.344e-01	2.513e-01	3.320	0.000986 ***
horsepower:weight	5.529e-05	5.227e-06	10.577	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.931 on 383 degrees of freedom
Multiple R-squared: 0.8618, Adjusted R-squared: 0.859
F-statistic: 298.6 on 8 and 383 DF, p-value: < 2.2e-16

[6]: *#hidden test cases*

- (c) Add another predictor that is the square of horsepower to the existing set of predictors (including the interaction term in (b)) and define a fit

```
[7]: #lm.fit = ?
# your code here
lm.fit = lm(mpg ~ cylinders + displacement + horsepower + weight + acceleration +
  + year + origin + horsepower:weight + horsepower^2
  , data=Auto)
summary(lm.fit)
```

Call:

```
lm(formula = mpg ~ cylinders + displacement + horsepower + weight +
  acceleration + year + origin + horsepower:weight + horsepower^2,
  data = Auto)
```

Residuals:

Min	1Q	Median	3Q	Max
-8.589	-1.617	-0.184	1.541	12.001

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.876e+00	4.511e+00	0.638	0.524147
cylinders	-2.955e-02	2.881e-01	-0.103	0.918363
displacement	5.950e-03	6.750e-03	0.881	0.378610
horsepower	-2.313e-01	2.363e-02	-9.791	< 2e-16 ***
weight	-1.121e-02	7.285e-04	-15.393	< 2e-16 ***
acceleration	-9.019e-02	8.855e-02	-1.019	0.309081

```

year                7.695e-01  4.494e-02  17.124  < 2e-16 ***
origin              8.344e-01  2.513e-01   3.320  0.000986 ***
horsepower:weight   5.529e-05  5.227e-06  10.577  < 2e-16 ***
---

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.931 on 383 degrees of freedom

Multiple R-squared: 0.8618, Adjusted R-squared: 0.859

F-statistic: 298.6 on 8 and 383 DF, p-value: < 2.2e-16

```
[8]: #hidden test cases
```

- (d) Create a binary variable, `mpg01`, that contains a 1 if `mpg` contains a value above its median, and a 0 if `mpg` contains a value below its median. You can compute the median using the `median()` function.

```
[9]: #mpg01 = ?
# your code here
mpg01 = ifelse(Auto$mpg > median(mpg), 1, 0)
mpg01
```

```

1. 0 2. 0 3. 0 4. 0 5. 0 6. 0 7. 0 8. 0 9. 0 10. 0 11. 0 12. 0 13. 0 14. 0 15. 1 16. 0 17. 0 18. 0 19. 1 20. 1
21. 1 22. 1 23. 1 24. 1 25. 0 26. 0 27. 0 28. 0 29. 0 30. 1 31. 1 32. 1 33. 0 34. 0 35. 0 36. 0 37. 0 38. 0
39. 0 40. 0 41. 0 42. 0 43. 0 44. 0 45. 0 46. 0 47. 0 48. 0 49. 1 50. 1 51. 1 52. 1 53. 1 54. 1 55. 1 56. 1
57. 1 58. 1 59. 1 60. 0 61. 0 62. 0 63. 0 64. 0 65. 0 66. 0 67. 0 68. 0 69. 0 70. 0 71. 0 72. 0 73. 0 74. 0
75. 0 76. 0 77. 0 78. 0 79. 1 80. 0 81. 1 82. 1 83. 1 84. 1 85. 0 86. 0 87. 0 88. 0 89. 0 90. 0 91. 0 92. 0
93. 0 94. 0 95. 0 96. 0 97. 0 98. 0 99. 0 100. 0 101. 1 102. 1 103. 0 104. 0 105. 0 106. 0 107. 0 108. 0
109. 0 110. 0 111. 0 112. 0 113. 0 114. 1 115. 0 116. 0 117. 1 118. 1 119. 0 120. 0 121. 0 122. 1 123. 0
124. 0 125. 0 126. 0 127. 0 128. 1 129. 1 130. 1 131. 1 132. 0 133. 0 134. 0 135. 0 136. 0 137. 0 138. 0
139. 0 140. 1 141. 1 142. 1 143. 1 144. 1 145. 1 146. 1 147. 1 148. 1 149. 1 150. 1 151. 0 152. 0 153. 0
154. 0 155. 0 156. 0 157. 0 158. 0 159. 0 160. 0 161. 0 162. 0 163. 0 164. 0 165. 0 166. 1 167. 1 168. 0
169. 1 170. 1 171. 1 172. 1 173. 0 174. 1 175. 0 176. 1 177. 1 178. 0 179. 1 180. 1 181. 1 182. 1 183. 1
184. 1 185. 1 186. 0 187. 0 188. 0 189. 0 190. 0 191. 0 192. 1 193. 0 194. 1 195. 1 196. 1 197. 1 198. 0
199. 0 200. 0 201. 0 202. 1 203. 1 204. 1 205. 1 206. 0 207. 0 208. 0 209. 0 210. 0 211. 0 212. 0 213. 0
214. 0 215. 1 216. 1 217. 1 218. 1 219. 1 220. 0 221. 0 222. 0 223. 0 224. 0 225. 0 226. 0 227. 0 228. 0
229. 0 230. 0 231. 0 232. 1 233. 1 234. 1 235. 1 236. 1 237. 1 238. 1 239. 1 240. 0 241. 0 242. 0 243. 1
244. 1 245. 1 246. 1 247. 1 248. 0 249. 0 250. 0 251. 0 252. 0 253. 0 254. 1 255. 0 256. 0 257. 0 258. 0
259. 0 260. 0 261. 0 262. 0 263. 0 264. 0 265. 1 266. 1 267. 1 268. 1 269. 0 270. 1 271. 1 272. 1 273. 0
274. 0 275. 0 276. 0 277. 1 278. 1 279. 0 280. 0 281. 0 282. 0 283. 0 284. 0 285. 0 286. 0 287. 0 288. 0
289. 0 290. 0 291. 0 292. 1 293. 1 294. 1 295. 1 296. 1 297. 1 298. 1 299. 1 300. 1 301. 1 302. 1 303. 1
304. 1 305. 1 306. 1 307. 1 308. 1 309. 1 310. 1 311. 1 312. 1 313. 1 314. 1 315. 0 316. 1 317. 1 318. 1
319. 1 320. 1 321. 1 322. 1 323. 1 324. 1 325. 1 326. 1 327. 1 328. 1 329. 1 330. 1 331. 1 332. 1 333. 1
334. 1 335. 1 336. 1 337. 1 338. 1 339. 1 340. 1 341. 1 342. 1 343. 1 344. 1 345. 1 346. 1 347. 1 348. 1
349. 1 350. 1 351. 1 352. 1 353. 1 354. 1 355. 1 356. 1 357. 1 358. 1 359. 0 360. 1 361. 0 362. 0 363. 1
364. 1 365. 1 366. 1 367. 1 368. 1 369. 1 370. 1 371. 1 372. 1 373. 1 374. 1 375. 1 376. 1 377. 1 378. 1
379. 1 380. 1 381. 1 382. 1 383. 1 384. 0 385. 1 386. 1 387. 1 388. 1 389. 1 390. 1 391. 1 392. 1

```

```
[10]: #hidden test cases
```

```
[11]: #Adding mpg01 to the data frame and removing mpg
Auto$mpg01 = mpg01
Auto$mpg = NULL

#splitting data into training and test data sets
n = dim(Auto)[1]
inds.train = sample(1:n,3*n/4)
Auto.train = Auto[inds.train,]
inds.test = (1:n)[-inds.train]
Auto.test = Auto[inds.test,]
```

```
[12]: head(Auto.train)
```

		cylinders	displacement	horsepower	weight	acceleration	year	origin	name
		<int>	<dbl>	<int>	<int>	<dbl>	<int>	<int>	<fct>
A data.frame: 6 × 9	160	8	351	148	4657	13.5	75	1	ford
	231	8	350	170	4165	11.4	77	1	chev
	345	4	86	64	1875	16.4	81	1	plym
	179	4	120	88	2957	17.0	75	2	peug
	260	6	200	85	3070	16.7	78	1	merc
	144	4	97	78	2300	14.5	74	2	opel

- (e) Perform LDA on the training data in order to predict mpg01 using the variables that seemed most associated (cylinders, displacement, weight) with mpg01. What is the test error of the model obtained(accuracy)?

```
[13]: #accuracy = ?
lda.fit = lda(mpg01 ~ cylinders + displacement + weight, data=Auto.train,
↳subset=inds.train)

lda.accuracy = function(model=lda.fit){

  # your code here
  lda.pred = predict(model, Auto.test)
  lda.class = lda.pred$class
  return (mean(lda.class == Auto.test$mpg01))
}

accuracy = lda.accuracy()
```

```
[14]: lda.accuracy()
```

0.877551020408163

```
[15]: #hidden test cases
```

- (f) Perform QDA similar to (e) and obtain the accuracy of the model??

```
[16]: #accuracy = ?

qda.fit = qda(mpg01 ~ cylinders + displacement + weight, data=Auto.train,
↳subset=inds.train)

qda.accuracy = function(model=qda.fit){

  # your code here
  qda.pred = predict(model, Auto.test)
  qda.class = qda.pred$class
  return (mean(qda.class == Auto.test$mpg01))
}

accuracy = qda.accuracy()
```

```
[17]: qda.accuracy()
```

```
0.86734693877551
```

```
[18]: #hidden test cases
```

(g) Perform NaiveBayes similar to (e) and obtain the accuracy of the model??

```
[19]: #accuracy = ?

nb.fit = naiveBayes(mpg01 ~ cylinders + displacement + weight, data=Auto.train,
↳subset=inds.train)

nb.accuracy = function(model=nb.fit){

  # your code here
  nb.class = predict(model, Auto.test)

  return(mean(nb.class == Auto.test$mpg01))
}

accuracy = nb.accuracy()
```

```
[20]: nb.accuracy()
```

```
0.86734693877551
```

```
[21]: #hidden test cases
```

```
[ ]:
```